
Deep Learning and Applications

Theja Tulabandhula

Today's Outline

- Feedforward Neural Nets
- Convolutional Neural Nets
 - Convolution
 - Pooling

Feedforward Neural Network

- Linear model $f(x, W, b) = Wx + b$
- A feedforward neural network model will include nonlinearities
- Two layer model
 - $f(x, W_1, b_1, W_2, b_2) = W_2 \max(0, W_1 x + b_1) + b_2$
 - Say x is d dimensional
- The number of hidden nodes is q
- The number of labels is p
- The notion of layer is for vectorizing/is conceptual

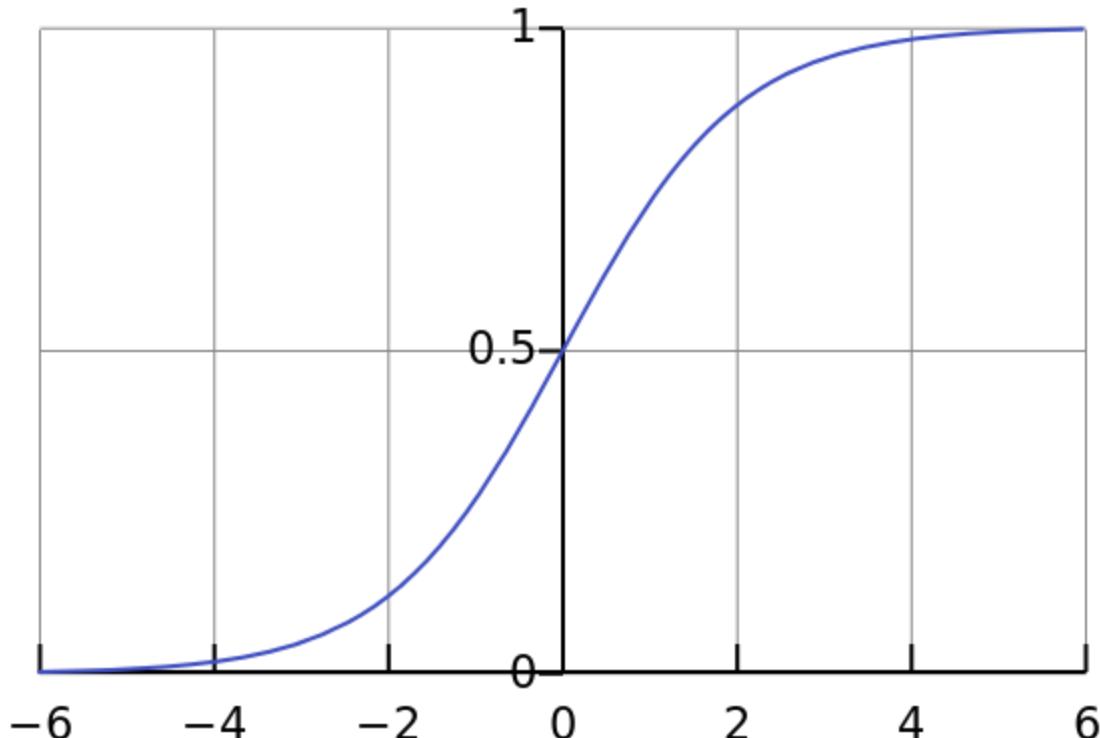
Nonlinearities (I)

Name	Formula	Year
none	$y = x$	-
sigmoid	$y = \frac{1}{1+e^{-x}}$	1986
tanh	$y = \frac{e^{2x}-1}{e^{2x}+1}$	1986
ReLU	$y = \max(x, 0)$	2010
(centered) SoftPlus	$y = \ln(e^x + 1) - \ln 2$	2011
LReLU	$y = \max(x, \alpha x), \alpha \approx 0.01$	2011
maxout	$y = \max(W_1x + b_1, W_2x + b_2)$	2013
APL	$y = \max(x, 0) + \sum_{s=1}^S a_i^s \max(0, -x + b_i^s)$	2014
VLReLU	$y = \max(x, \alpha x), \alpha \in 0.1, 0.5$	2014
RReLU	$y = \max(x, \alpha x), \alpha = \text{random}(0.1, 0.5)$	2015
PReLU	$y = \max(x, \alpha x), \alpha \text{ is learnable}$	2015
ELU	$y = x, \text{ if } x \geq 0, \text{ else } \alpha(e^x - 1)$	2015

- How to pick the nonlinearity/activation function?

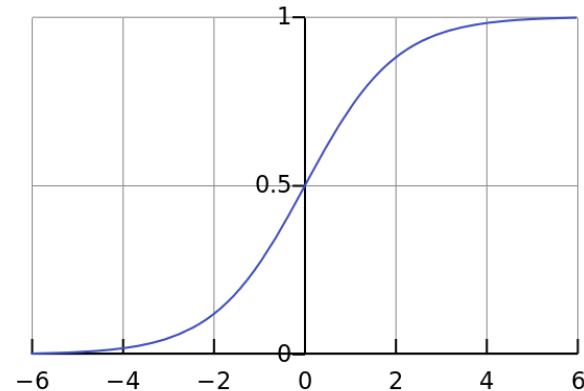
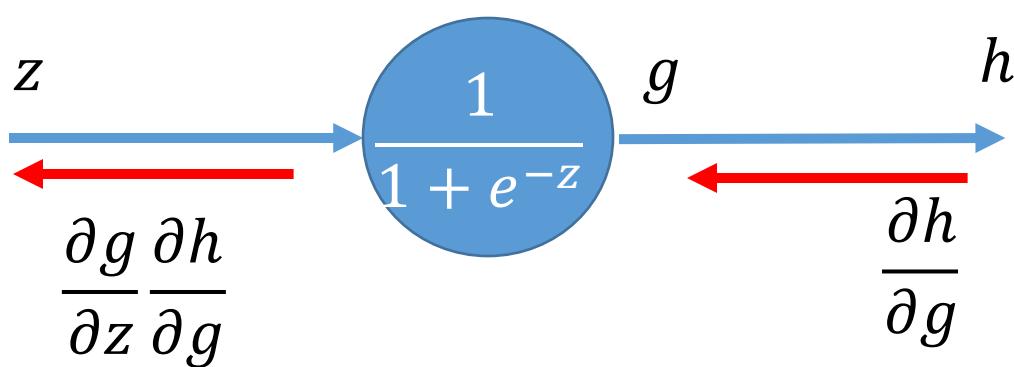
Nonlinearities (II)

- Sigmoid
 - Is a map whose range is $[0,1]$



Nonlinearities (III)

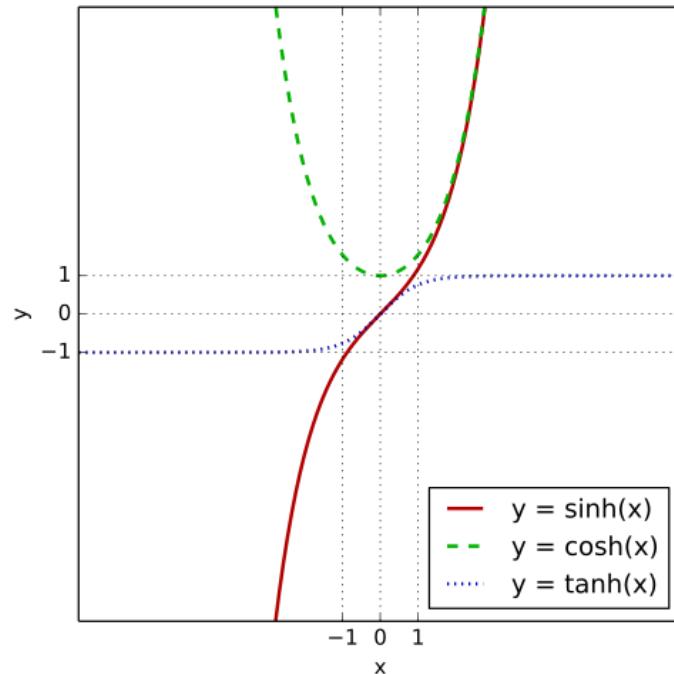
- Saturated node/neuron makes gradients vanish



- Not zero-centered
 - Empirically may lead to slower convergence

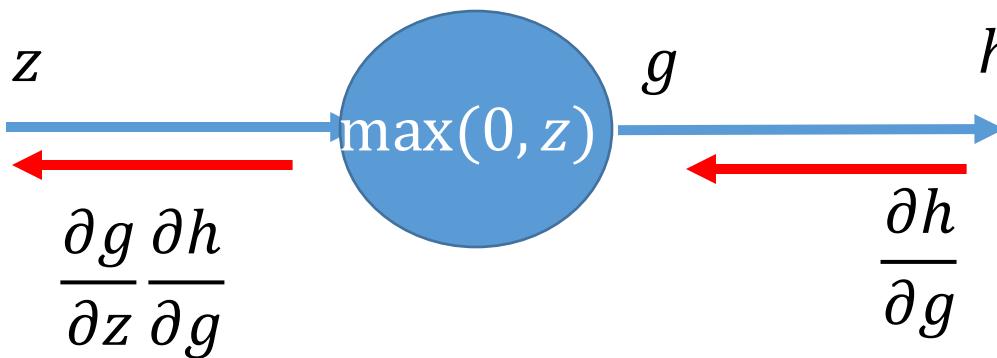
Nonlinearities (IV)

- `tanh()` addresses the zero-centering problem. So will typically give better results
- Still gradients vanish

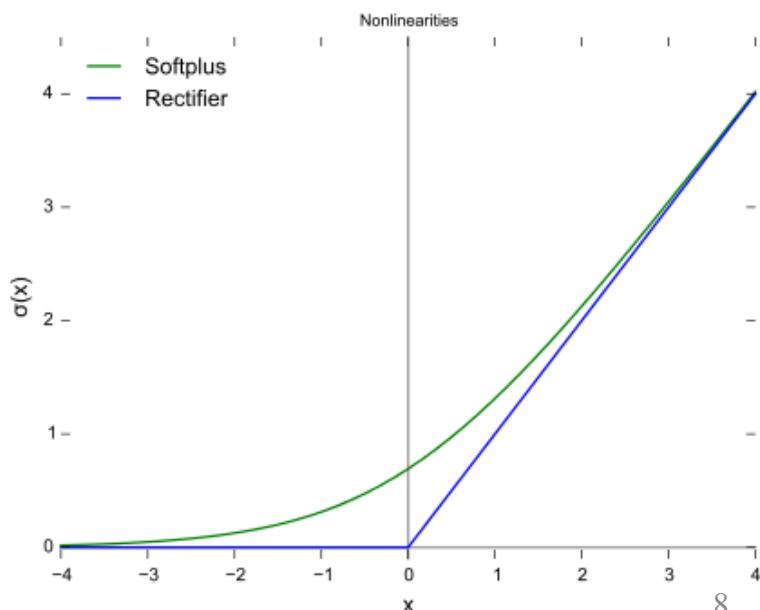


Nonlinearities (V)

- ReLU (2012 Krizhevsky et al.)
- No vanishing gradient on the positive side
- Empirically observed to be very good
- Initialization/high learning rate may lead to permanently dead ReLUs (diagnosable)

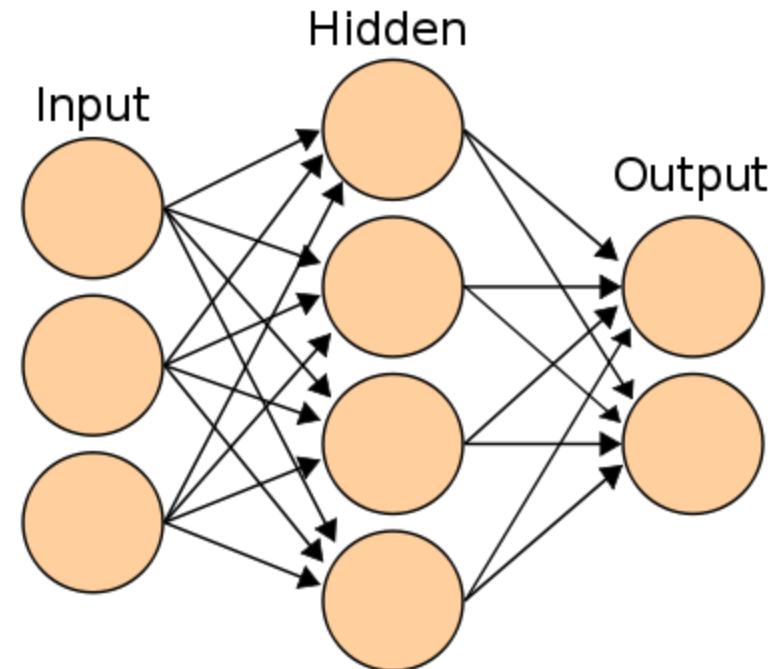


Is a gradient gate!



Feedforward Neural Net

- Lets focus on a 2-layer net
- Layers
 - Input
 - Hidden
 - Output
- Node
- Nonlinearity
 - Activation



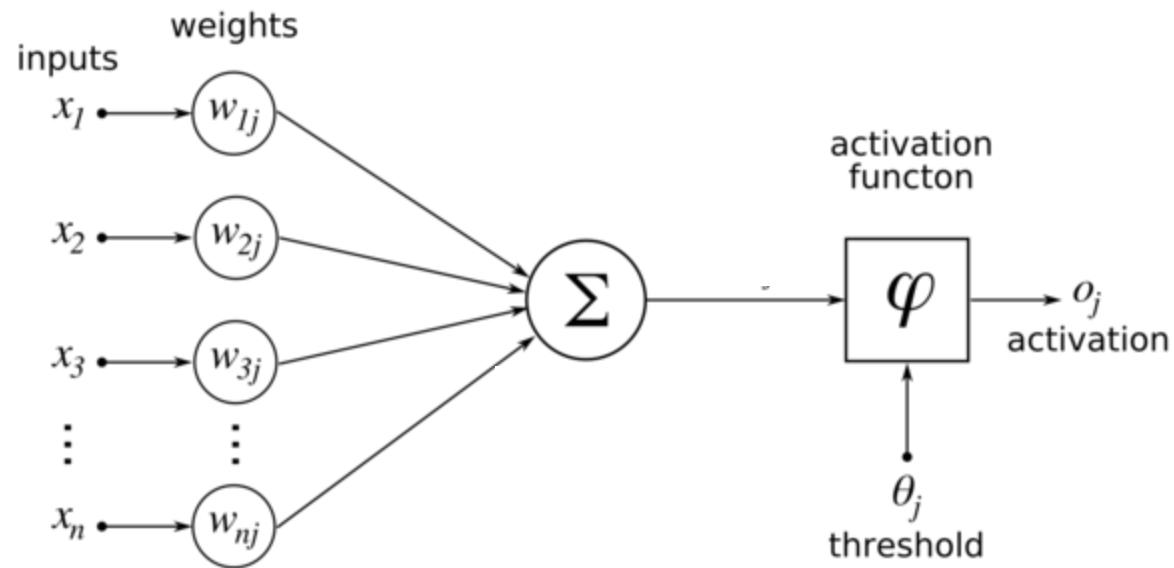
$$f(x, W_1, b_1, W_2, b_2) = W_2 \max(0, W_1 x + b_1) + b_2$$

Neuron (I)

- Historical
- Let $f(x) = w \cdot x + b$
- Perceptron from 1957: $h(x) = \begin{cases} 0, & f(x) < 0 \\ 1, & \text{otherwise} \end{cases}$
- Update rule was $w_{k+1} = w_k + \alpha(y - h(x))x$ similar to gradient update rules we see today
- Passing the score through a sigmoid was likened to how a neuron fires
 - Firing rate = $\frac{1}{1+e^{-yf(x)}}$

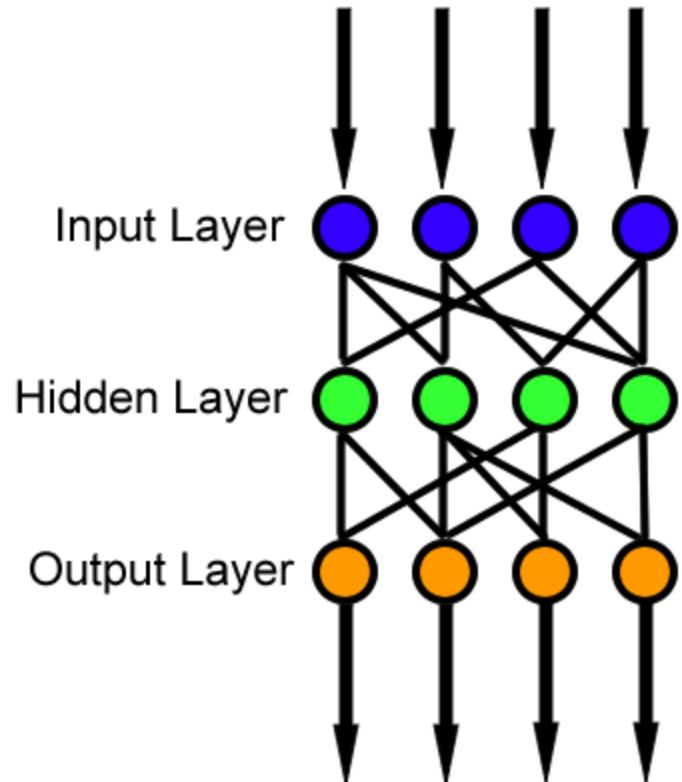
Neuron (II)

- Get a vector x_i and transform it to a score vector by passing through a sequence of hidden layers
- Each hidden layer has neurons
- Each neuron is fully connected to previous layer



Feedforward Net: Two Layer Model

- Number of layers is the number of W, b pairs
- Some questions to think about:
 - How to pick the number of layers?
 - How to pick the number of hidden units in each layer?



Feedforward Net and Backprop

- Choose a mini-batch (sample) of size B
- Forward propagate through the computation graph
 - Compute losses $L_{i_1}, L_{i_2}, \dots L_{i_B}$ and $R(W_1, b_1, W_2, b_2)$
 - Get loss L for the batch
- Backprop to compute gradients with respect to W_1, b_1, W_2 and b_2
- Update parameters W_1, b_1, W_2 and b_2
 - In the direction of the negative gradient

Feedforward Net in Python

```
# Feedforward neural net model

# Start with an initial set of parameters randomly
h = 100 # size of hidden layer
W = 0.01 * np.random.randn(D,h)
b = np.zeros((1,h))
W2 = 0.01 * np.random.randn(h,K)
b2 = np.zeros((1,K))

# Initial values from hyperparameter
reg = 1e-3 # regularization strength

#For simplicity, we will not optimize this using grid search here.
```

Feedforward Net in Python

```
#Perform batch SGD using manual backprop

#For simplicity we will take the batch size to be the same as number of examples
num_examples = X.shape[0]

#Initial value for the Gradient Descent Parameter
step_size = 1e-0 #Also called learning rate

#For simplicity, we will not hand tune this algorithm parameter as well.

# gradient descent loop
for i in xrange(10000):

    # evaluate class scores, [N x K]
    hidden_layer = np.maximum(0, np.dot(X, W) + b) # note, ReLU activation
    scores = np.dot(hidden_layer, W2) + b2

    # compute the class probabilities
    exp_scores = np.exp(scores)
    probs = exp_scores / np.sum(exp_scores, axis=1, keepdims=True) # [N x K]

    # compute the loss: average cross-entropy loss and regularization
    corect_logprobs = -np.log(probs[range(num_examples),y])
    data_loss = np.sum(corect_logprobs)/num_examples
    reg_loss = 0.5*reg*np.sum(W*W) + 0.5*reg*np.sum(W2*W2)
    loss = data_loss + reg_loss
    if i % 1000 == 0:
        print "iteration %d: loss %f" % (i, loss)
```

Feedforward Net in Python

```
# compute the gradient on scores
dscores = probs
dscores[range(num_examples),y] -= 1
dscores /= num_examples

# backpropate the gradient to the parameters
# first backprop into parameters W2 and b2
dW2 = np.dot(hidden_layer.T, dscores)
db2 = np.sum(dscores, axis=0, keepdims=True)
# next backprop into hidden layer
dhidden = np.dot(dscores, W2.T)
# backprop the ReLU non-linearity
dhidden[hidden_layer <= 0] = 0
# finally into W,b
dW = np.dot(X.T, dhidden)
db = np.sum(dhidden, axis=0, keepdims=True)

# add regularization gradient contribution
dW2 += reg * W2
dW += reg * W

# perform a parameter update
W += -step_size * dW
b += -step_size * db
W2 += -step_size * dW2
b2 += -step_size * db2
```

Feedforward Net in Python

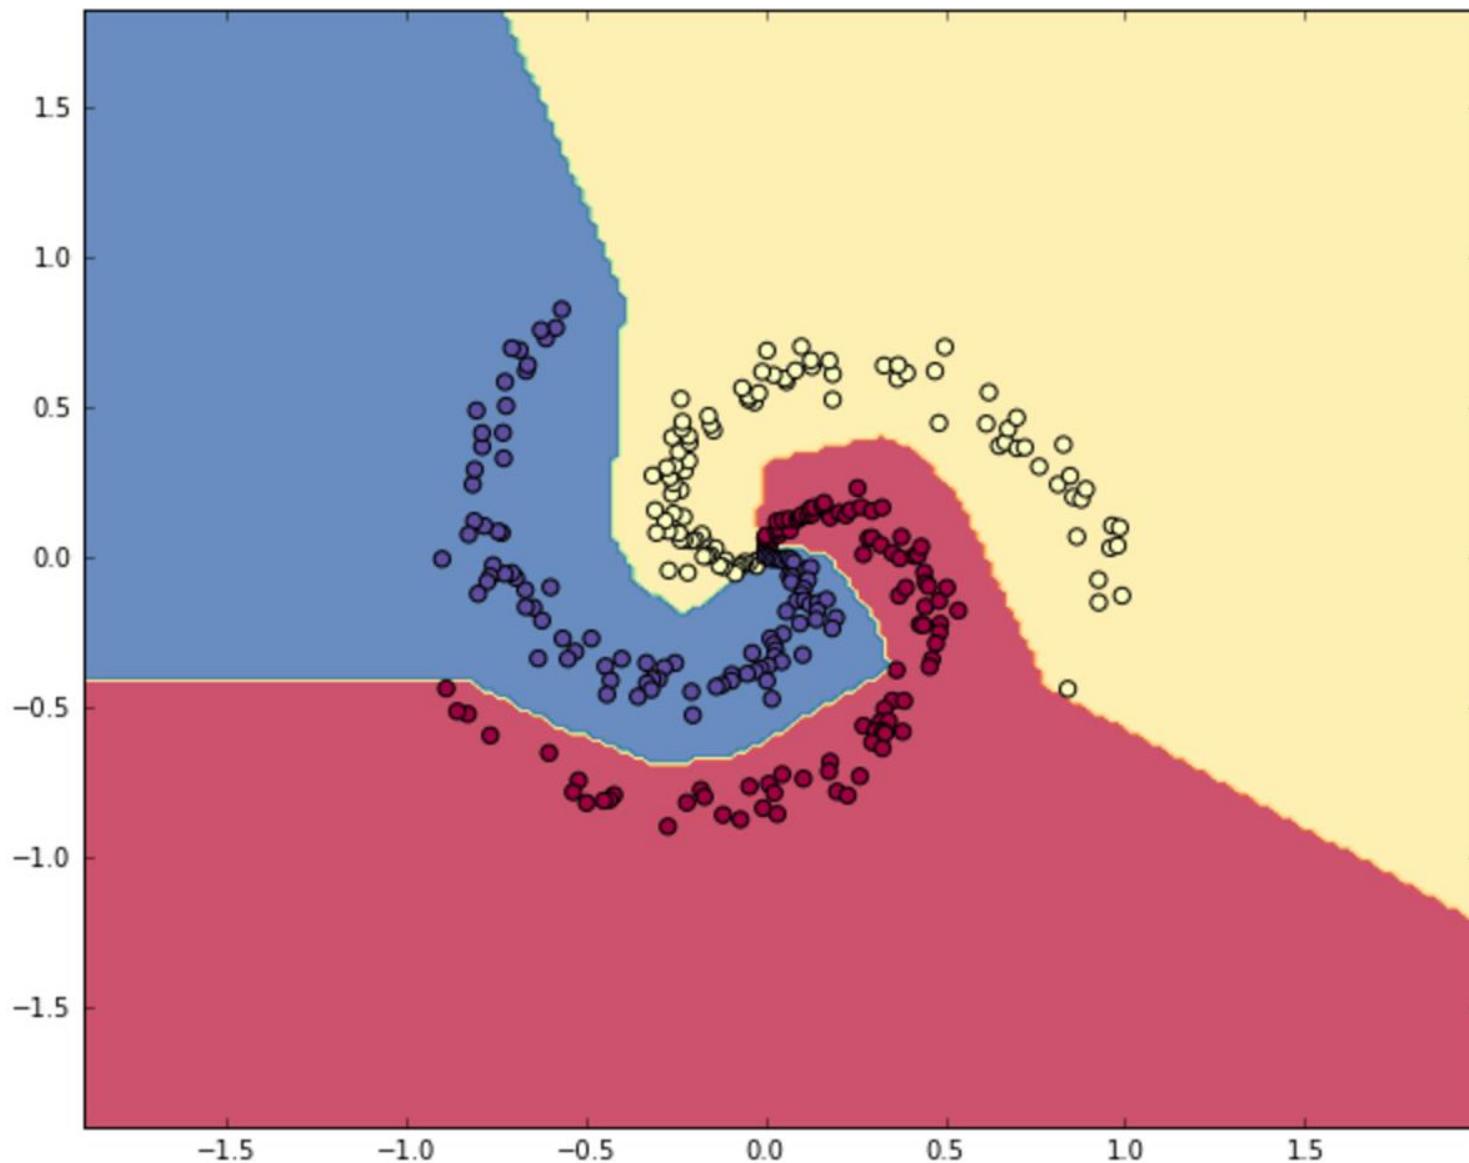
Post Training

```
# Post-training: evaluate test set accuracy

#For simplicity, we will use training data as proxy for test. Do not do this.
X_test = X
y_test = y

hidden_layer = np.maximum(0, np.dot(X_test, W) + b)
scores = np.dot(hidden_layer, W2) + b2
predicted_class = np.argmax(scores, axis=1)
print 'test accuracy: %.2f' % (np.mean(predicted_class == y_test))
```

Feedforward Net in Python



Questions?

Today's Outline

- Feedforward Neural Nets
- Convolutional Neural Nets
 - Convolution
 - Pooling

Convolutional Neural Network

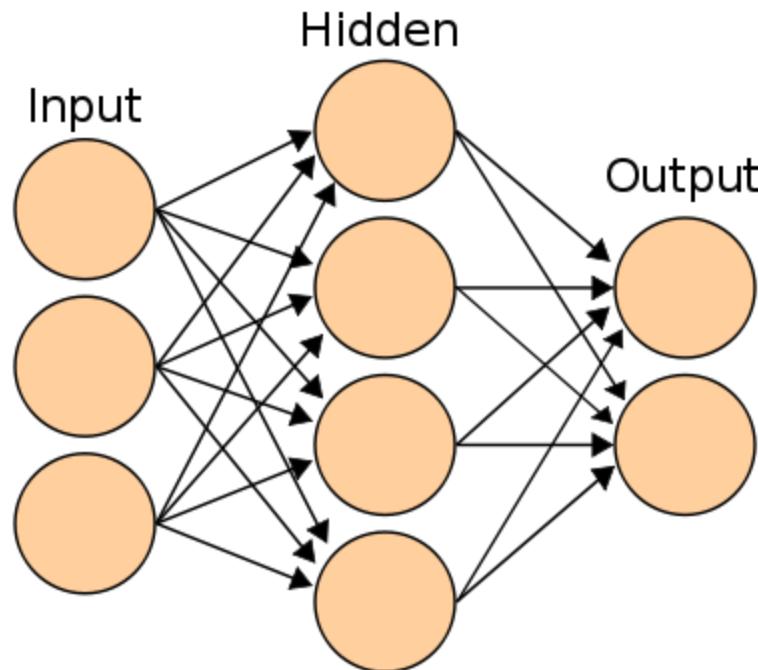
Similar to Feedforward NN

- Similar to feedforward neural networks
- Each neuron/node is associated with weights and a bias
- Node receives input
 - Performs dot product of vectors
 - Applies non-linearity
- The difference:
 - Number of parameters is reduced!

How? That is the content of this lecture!

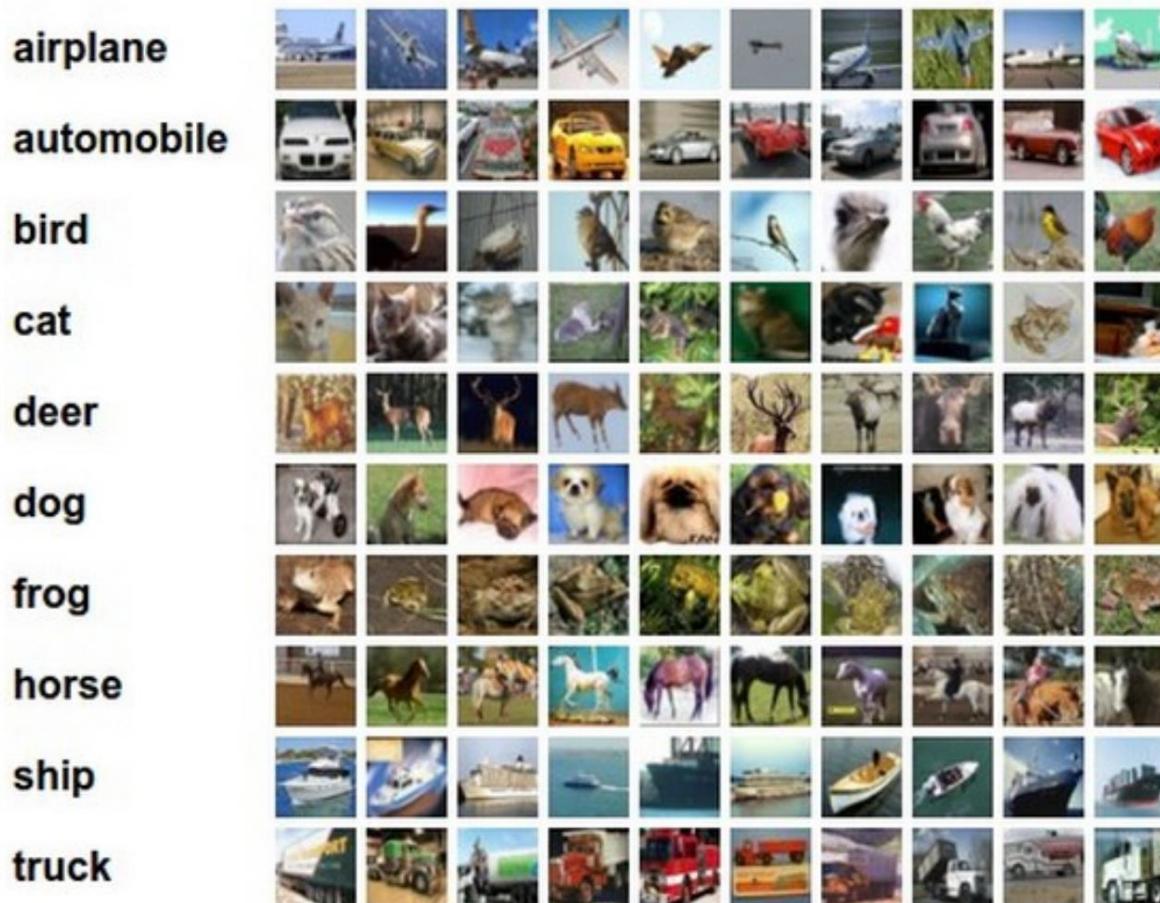
Towards CNNs (I)

- Feedforward net:
 - Can you visualize the connections for an arbitrary neuron here?



Towards CNNs (II)

- Consider the CIFAR-10 Dataset. Images are 32*32*3 in size



Towards CNNs (III)

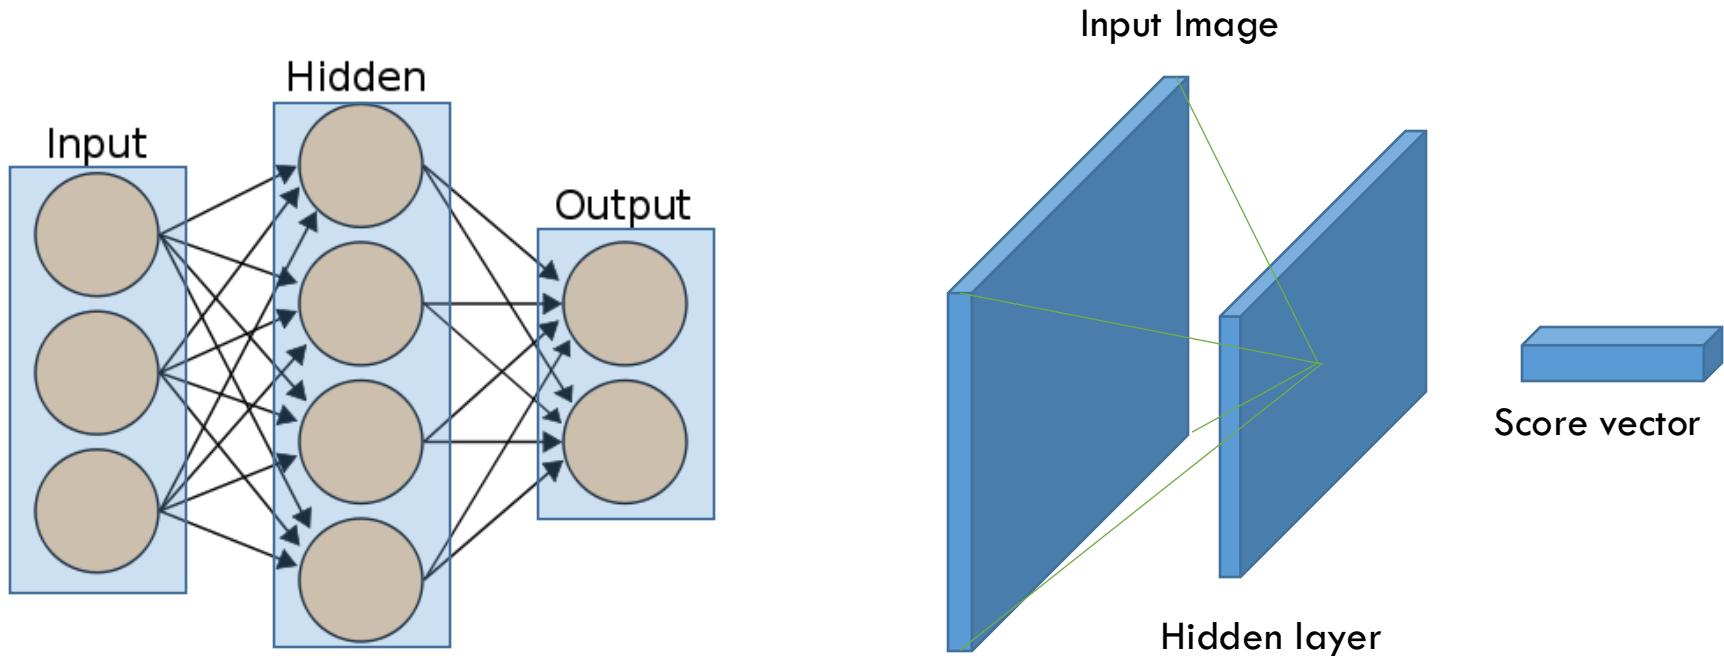
- First fully connected feedforward neuron would have $32*32*3$ weights associated with it (+1 bias parameter)
- What if the images were $1280*800*3$?
- Clearly, we also need many neurons in each hidden layer. This leads to explosion in the total number of parameters (or the dimension of W s and b s)

CNN Architecture

- We will look at it from layers point of view
- The new idea is that layers have **width** and **depth**!
 - (In contrast, Feedforward NN layers only had height)
 - (depth here does NOT correspond to number of layers of a network)

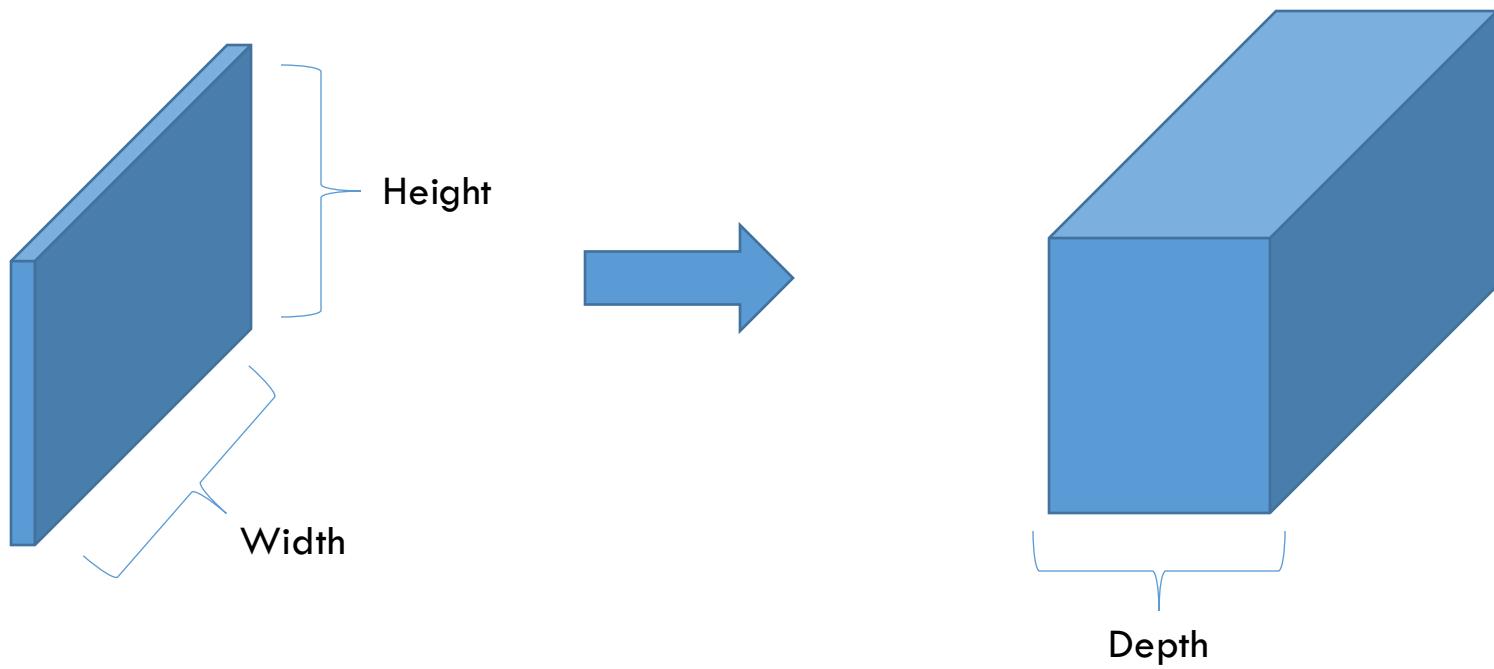
CNN Architecture

- View FFN layers as having width and height



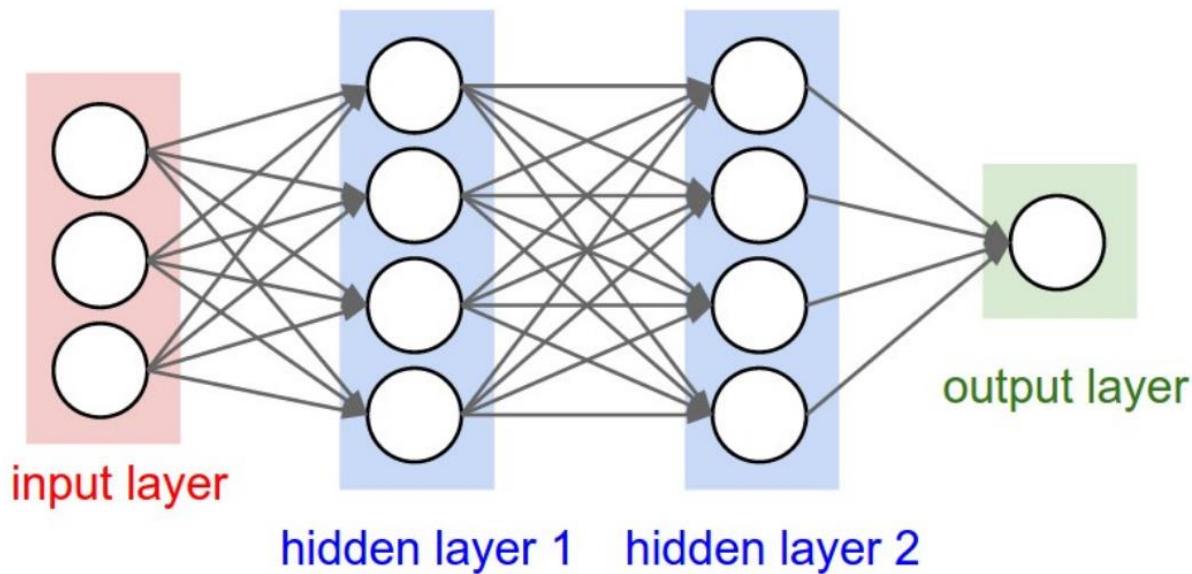
CNN Architecture

- The new idea is that CNN layers have **depth**!
 - (depth here does NOT correspond to number of layers of a network)



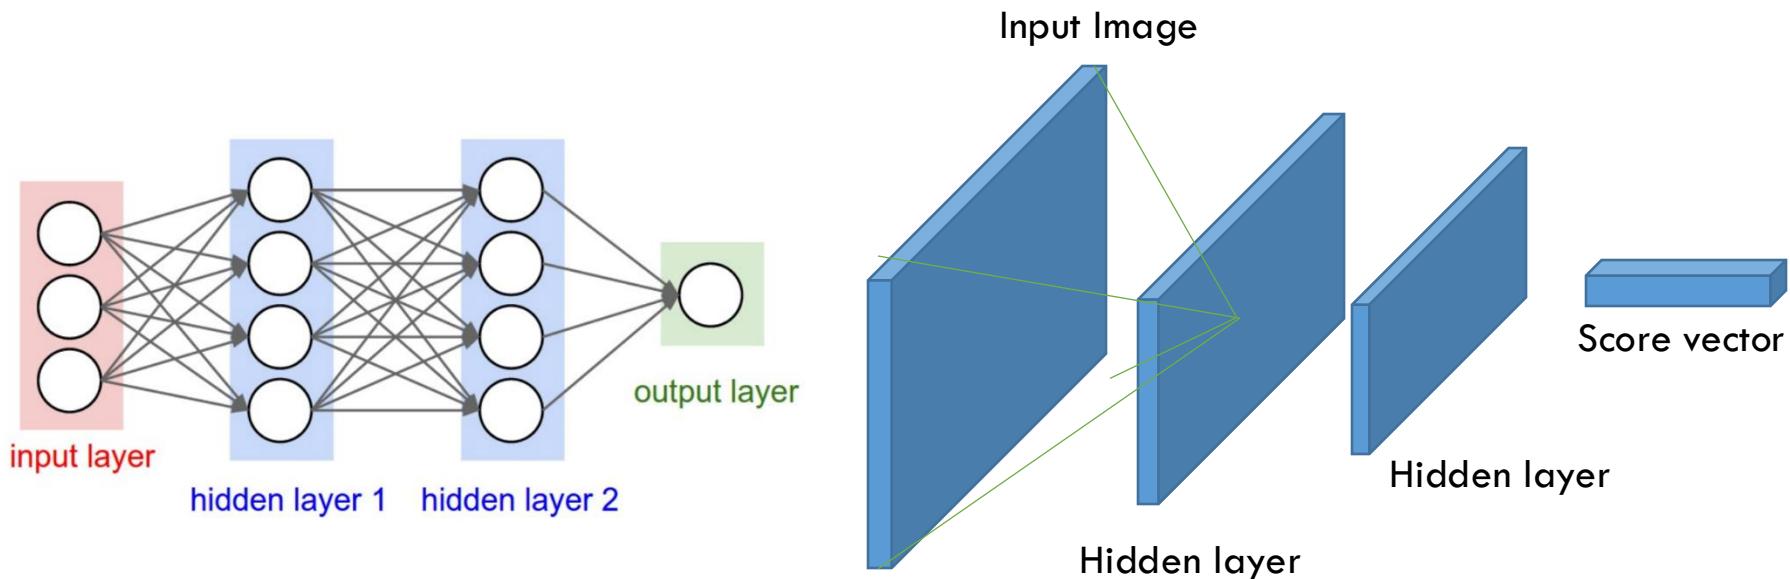
3D Volumes of Neurons

- Input has dimension $32*32*3$ (for CIFAR-10 dataset)
- Final output has dimension $1*1*10$ (10 classes)
- Previously,



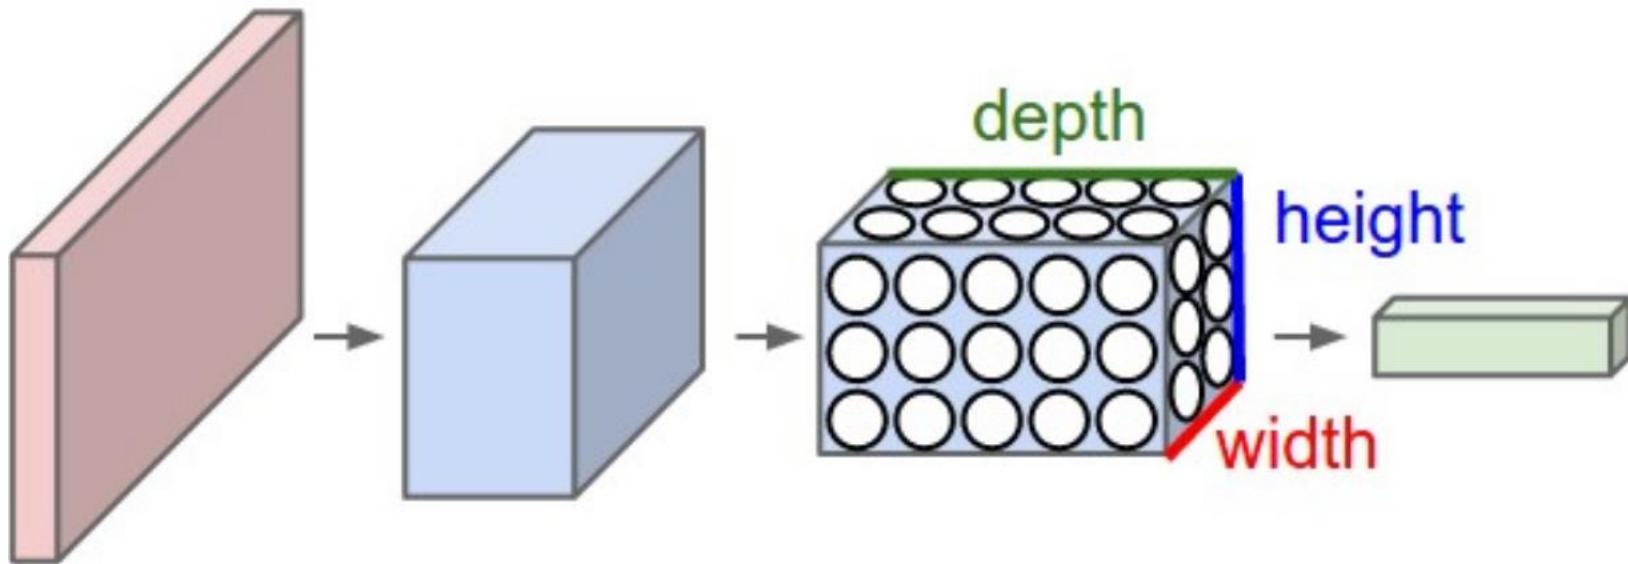
3D Volumes of Neurons

- Input has dimension $32*32*3$ (for CIFAR-10 dataset)
- Final output has dimension $1*1*10$ (10 classes)
- So assuming 2 hidden layers, previously we had,



3D Volumes of Neurons

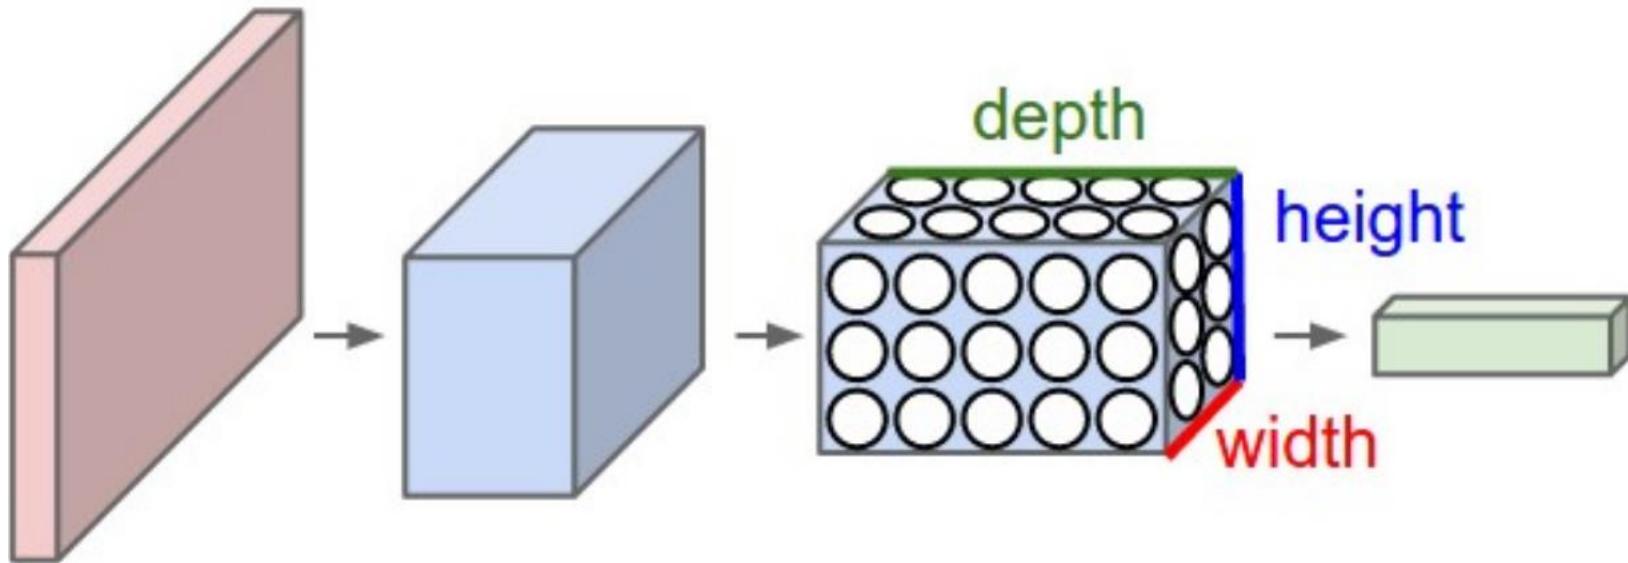
- Now,



- Each layer simply does this: transforms an input tensor (3D volume) to an output tensor using some function

3D Volumes of Neurons

- Now,



- Each layer simply does this: transforms an input tensor (3D volume) to an output tensor using some function

CNN Layers

- Three types
 - Convolutional Layer (CONV)
 - Pooling Layer (POOL)
 - Fully Connected Layer (same as Feedforward neural network, i.e., $1*1*\#$ Neurons is the layer's output tensor)
- Stack these in various ways

CNN Example Architecture

- Say our classification dataset is CIFAR-10
- Let the architecture be as follows:
 - INPUT -> CONV -> POOL -> FC
- INPUT:
 - This layer is nothing but $32 \times 32 \times 3$ in dimension (width*height*3 color channels)

CNN Example Architecture

- Say our classification dataset is CIFAR-10
- Let the architecture be as follows:
 - INPUT -> CONV -> POOL -> FC
- INPUT:
 - This layer is nothing but $32 \times 32 \times 3$ in dimension (width*height*3 color channels)
- CONV:
 - Neurons compute like regular feedforward neurons (sum the product of inputs with weights and add bias).
 - May output a different shaped tensor, say with dimension $32 \times 32 \times 12$

CNN Example Architecture

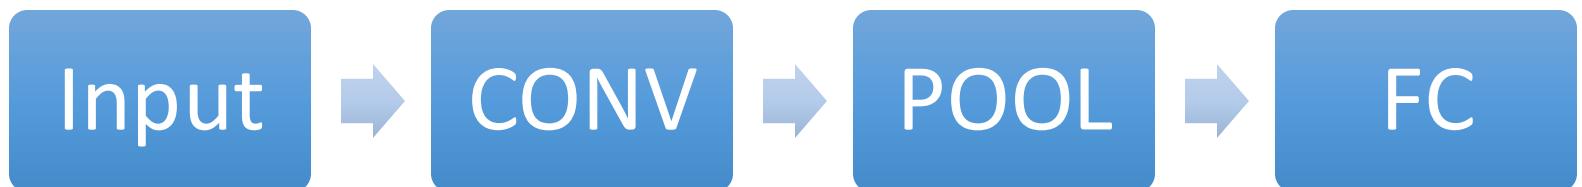
- POOL:
 - Performs a down-sampling in the spatial dimension
 - Outputs a tensor with the depth dimension the same as input
 - If input is $32*32*12$, then output could be $16*16*12$

CNN Example Architecture

- POOL:
 - Performs a down-sampling in the spatial dimension
 - Outputs a tensor with the depth dimension the same as input
 - If input is $32*32*12$, then output could be $16*16*12$
- FC:
 - This is the fully connected layer. Input can be any tensor (say $16*16*12$) but the output will have only one effective dimension ($1*1*10$ since this is the last layer and CIFAR-10 has 10 classes)

CNN Example Architecture

- So we went from pixels (32*32 RGB images) to scores (10 in number)



- Some layers have parameters (CONV and FC), other layers do not (POOL)
- Optimization of these parameters still for achieving scores consistent with image labels

The Convolution Layer (CONV)

- Layer's parameters correspond to a **set** of **filters**
- What is a filter?
 - A **linear function parameterized by a tensor**
 - Outputs a scalar
 - The parameter tensor is **learned** during training
- Example
 - First layer filter may be of dimension $3 \times 3 \times 3$
 - 3 pixels wide
 - 3 pixels high
 - 3 unit **filter-depth** for three color channels
- We slide (convolve) the filter **across the width and height** of the input volume and compute the scalar output to be passed into the nonlinearity

CONV: Sliding/Convolving

- We slide (convolve) the filter across the width and height of the input volume and compute the scalar output to be passed into the nonlinearity

1 <small>x1</small>	1 <small>x0</small>	1 <small>x1</small>	0	0
0 <small>x0</small>	1 <small>x1</small>	1 <small>x0</small>	1	0
0 <small>x1</small>	0 <small>x0</small>	1 <small>x1</small>	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

Convolved Feature

Also see <http://setosa.io/ev/image-kernels/>

¹Figure: http://deeplearning.stanford.edu/wiki/index.php/Feature_extraction_using_convolution

The Convolution Layer (CONV)

- Three things to notice
 - Filters are small along width and height
 - Same **filter-depth** as the input tensor (3D volume)
 - If the input is $x * y * z$, then filter could be $3 * 3 * z$
 - As we slide, we produce a **2D** activation map

The Convolution Layer (CONV)

- Three things to notice
 - Filters are small along width and height
 - Same **filter-depth** as the input tensor (3D volume)
 - If the input is $x * y * z$, then filter could be $3 * 3 * z$
 - As we slide, we produce a **2D** activation map
- Filters (i.e., filter parameters) will be learned during training that ‘detect’ certain visual features
 - Example:
 - Oriented edges, colors, etc. at the first layer
 - Specific patterns in higher layers

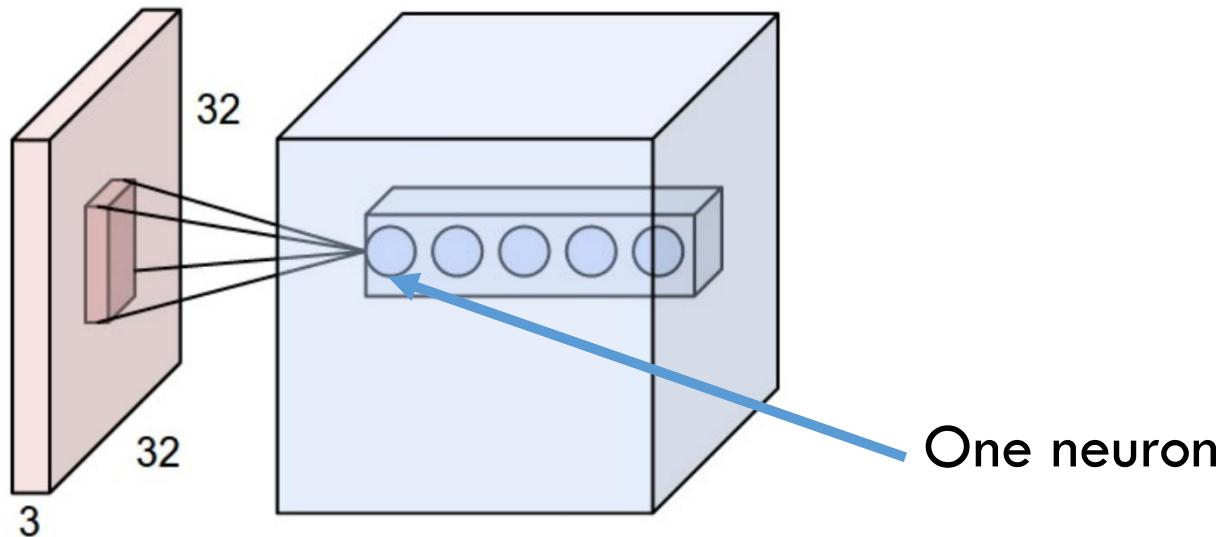
CONV: Filters

- Before we look at the patterns ...
- Lets now look at the neurons themselves
 - How are they connected?
 - How are they arranged?
 - How can we get reduced parameters?

CONV: Local Connectivity

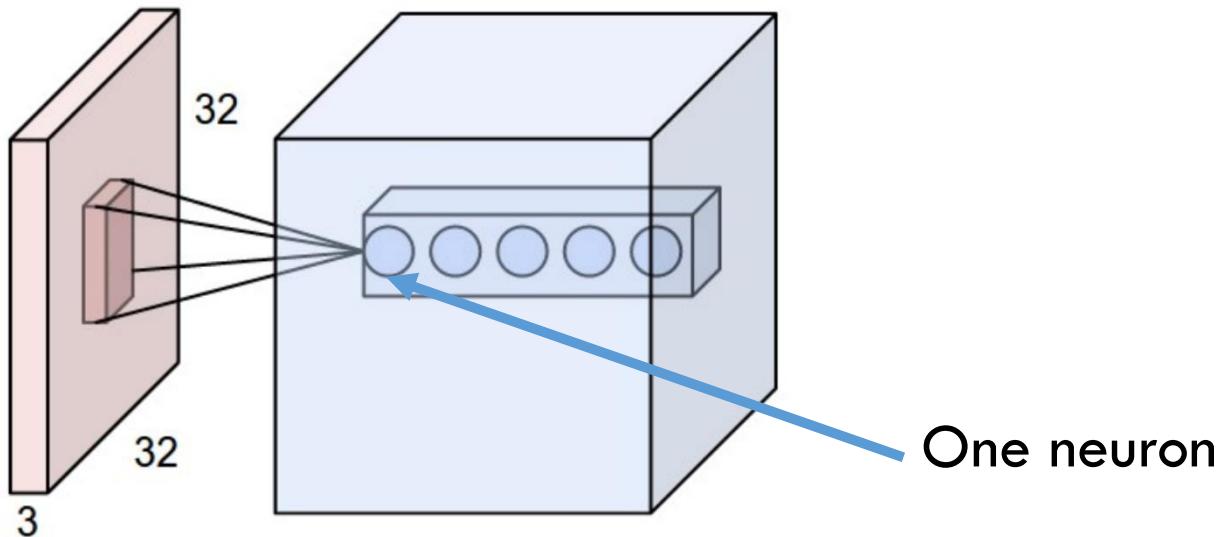
- Connect each neuron to a local (spatial) region of the input tensor
- Spatial extent of this connectivity is called **receptive field**
- Depth connectivity is the same as input depth

CONV: Local Connectivity



- Example: If input tensor is $32 \times 32 \times 3$ and filter is $3 \times 3 \times 3$ then
 - the number of weight parameters is 27, and
 - there is 1 bias parameter

CONV: Local Connectivity



- All 5 neurons are looking at the same spatial region
- Each neuron belongs to a different filter

CONV: Spatial Arrangement

- Back to layer point of view
- Size of output **tensor** depends on three numbers:
 - **Layer Depth**
 - Corresponds to the number of filters
 - **Stride** (how much the filter is moved spatial)
 - Example: If stride is 1, then filter is moved 1 pixel at a time
 - **Zero-padding**
 - Deals with boundaries (is usually 1 or 2)

CONV: Stride/Zero-pad

Stride = 1, Zero-padding = 0

The diagram illustrates a convolution operation. A blue arrow points from the text "Stride = 1, Zero-padding = 0" to a 5x5 input image matrix. The image matrix has alternating yellow and green cells. Red annotations show the result of element-wise multiplication between a 3x3 kernel and the input. The first row of the kernel is multiplied by the first row of the image, resulting in a row of 1s. The second row of the kernel is multiplied by the second row of the image, resulting in a row of 0s. The third row of the kernel is multiplied by the third row of the image, resulting in a row of 1s. The fourth row of the kernel is multiplied by the fourth row of the image, resulting in a row of 1s. The fifth row of the kernel is multiplied by the fifth row of the image, resulting in a row of 0s. The resulting 1x5 vector is then reduced to a single value of 4, which is highlighted in a pink box.

1 _{x1}	1 _{x0}	1 _{x1}	0	0
0 _{x0}	1 _{x1}	1 _{x0}	1	0
0 _{x1}	0 _{x0}	1 _{x1}	1	1
0	0	1	1	0
0	1	1	0	0

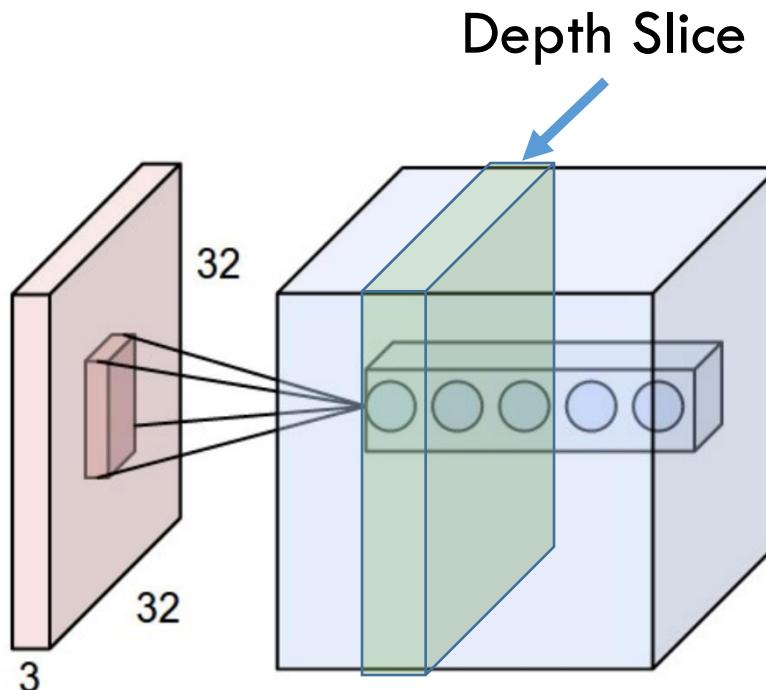
Image

4		

Convolved
Feature

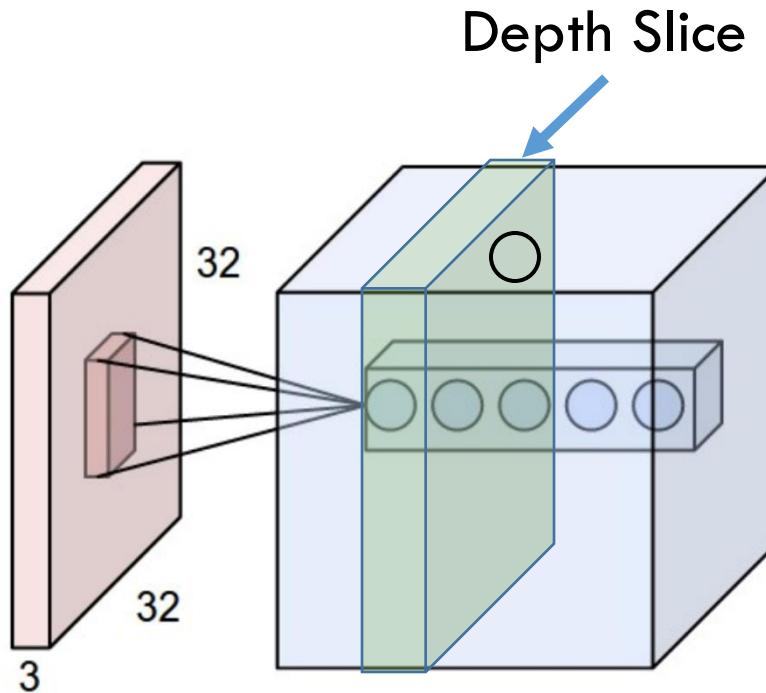
CONV: Parameter Sharing

- Key assumption: If a filter is useful for one region, it should also be useful for another region
- Denote a single 2D slice of depth of a layer as **depth slice**



CONV: Parameter Sharing

- Then, all neurons in each depth slice use the same weight and bias parameters!

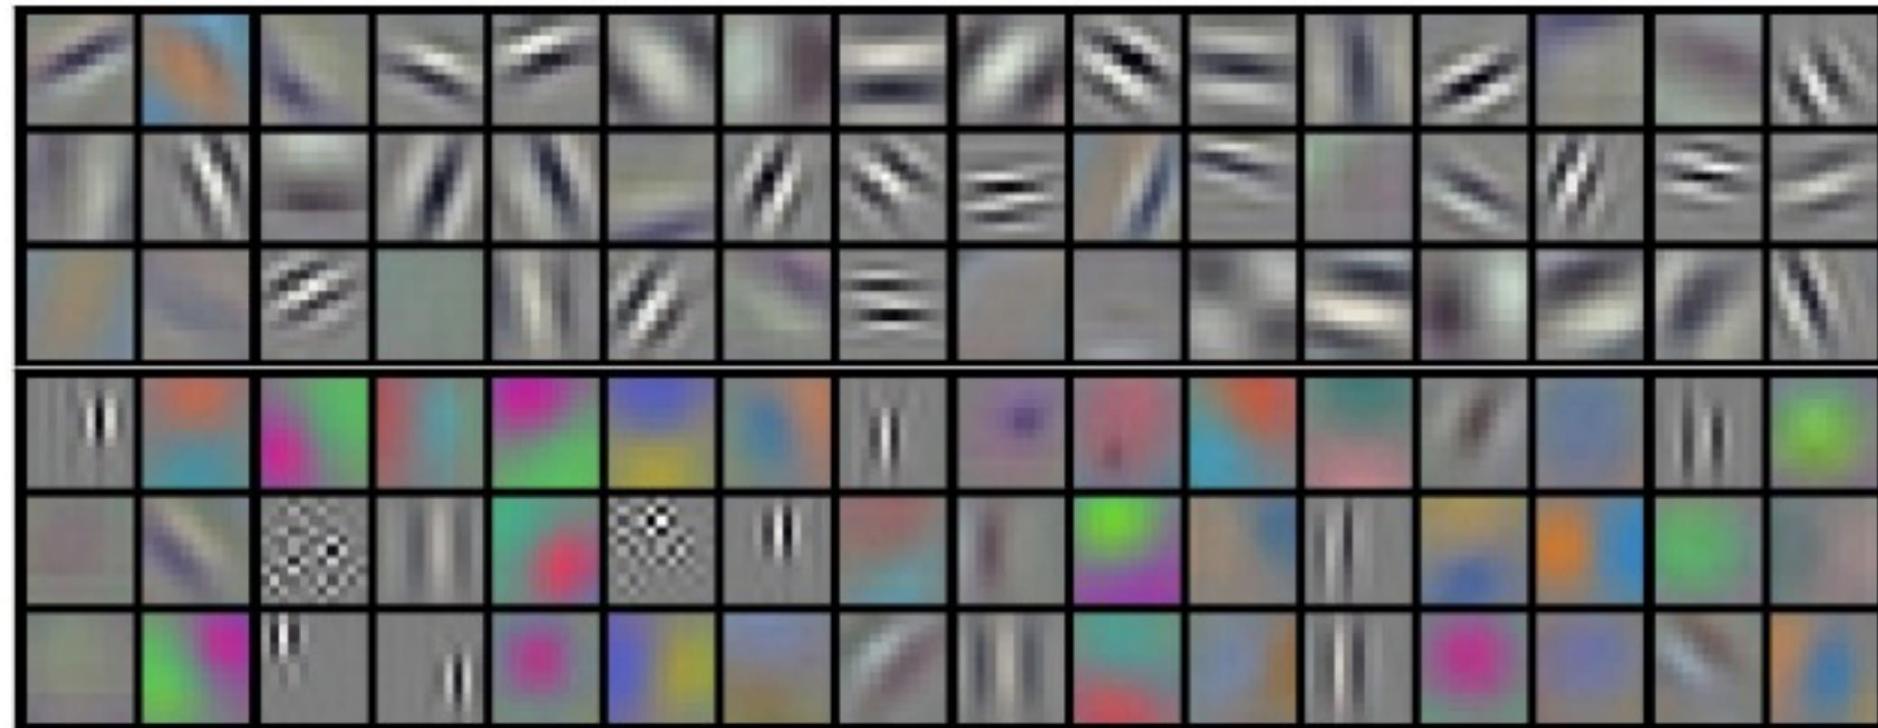


CONV: Parameter Sharing

- Number of parameters is reduced!
- Example:
 - Say the number of filters is M (= Layer Depth)
 - Then, this layer will have $M * (3 * 3 * 3 + 1)$ parameters
- Gradients will get added up across neurons of a depth slice

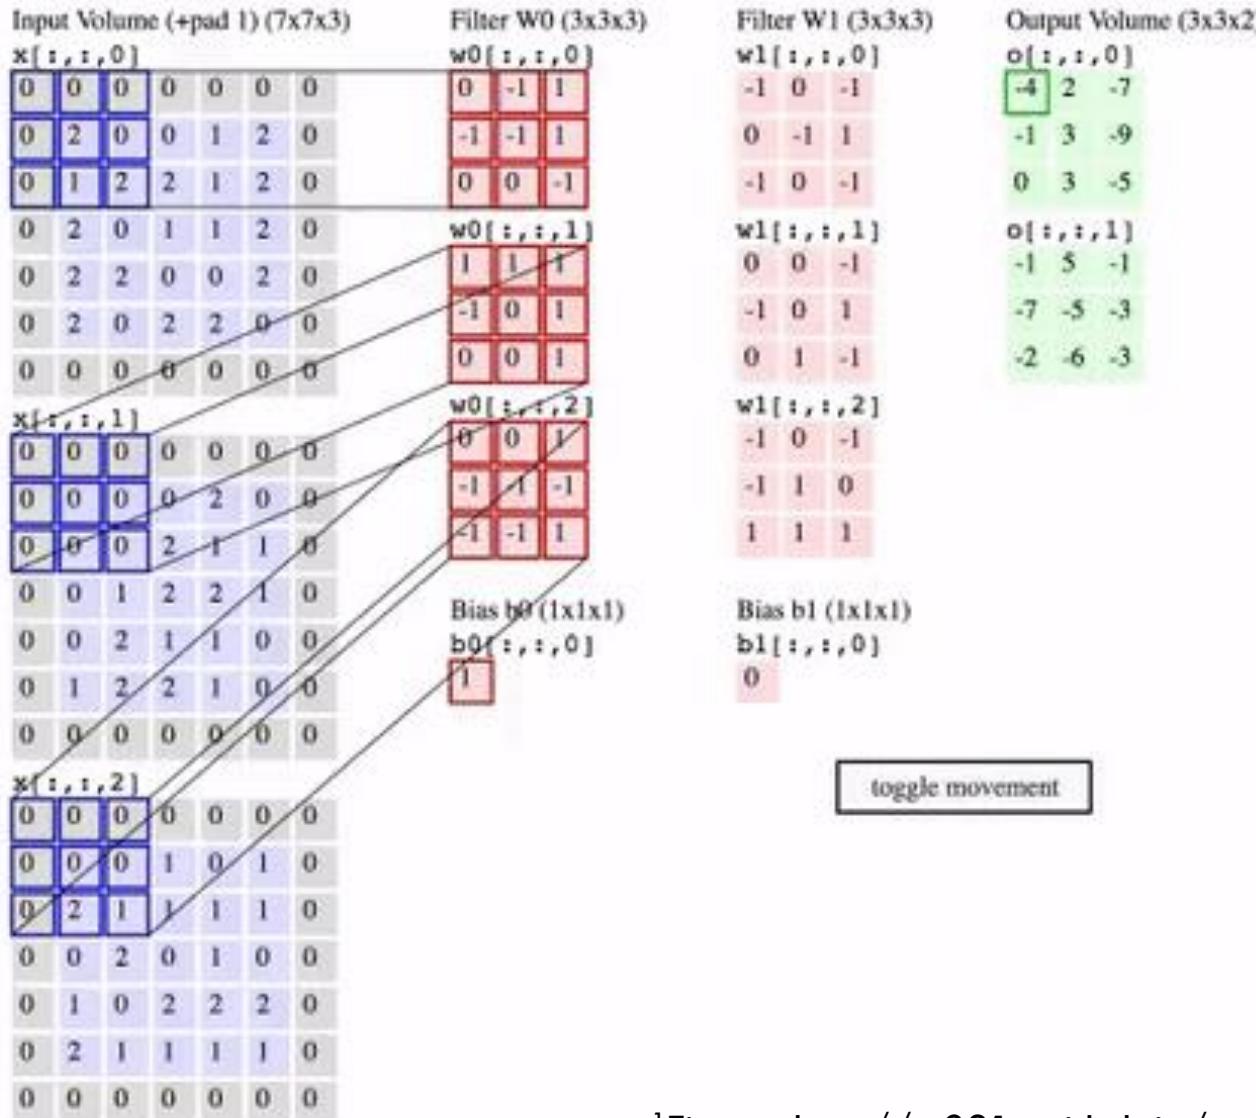
CONV: Parameter Sharing

- AlexNet's first layer has $11 \times 11 \times 3$ sized filters 96 in number. The filter weights are plotted below:



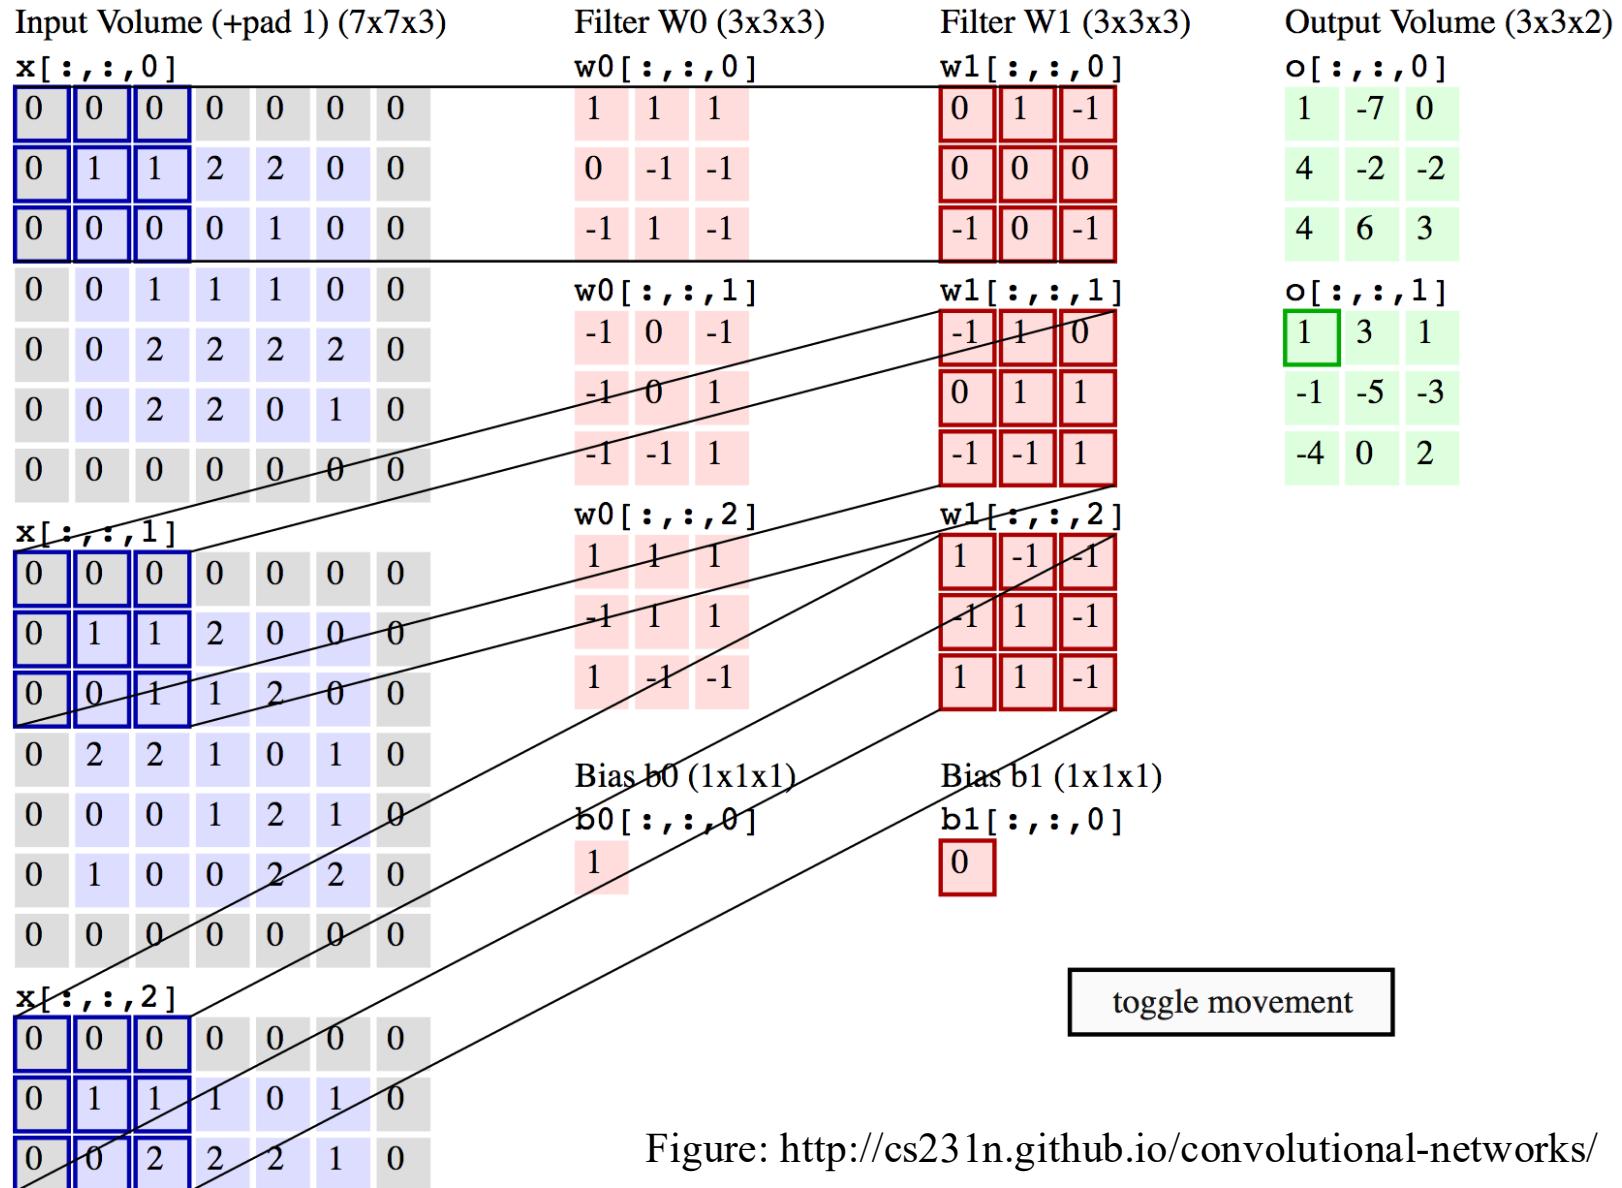
- Intuition: If capturing an edge is important, then important everywhere

Example: CONV Layer Computation



The name ‘convolution’ comes from the convolution operation in signal processing that is essentially a matrix matrix product.

Example: CONV Layer Computation

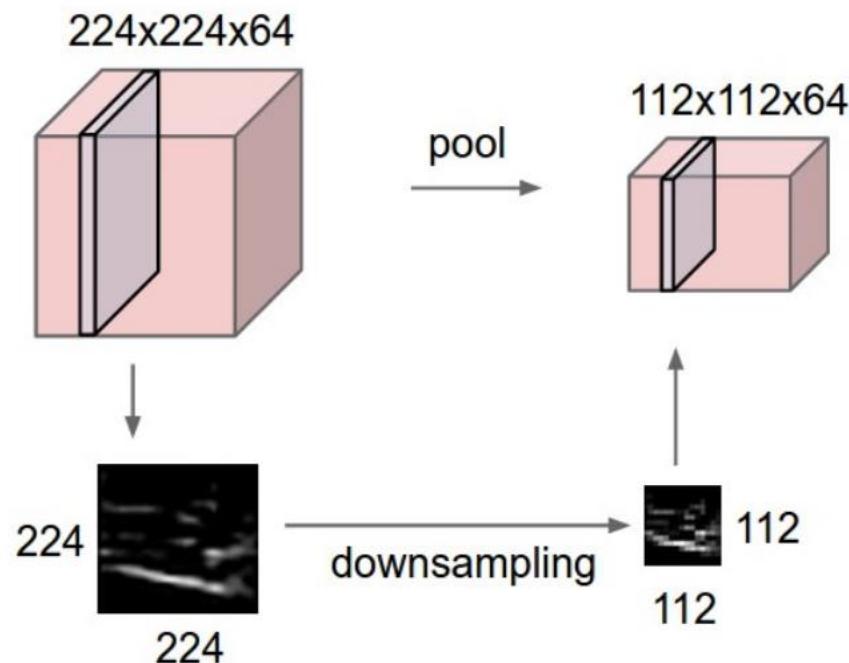


The Pooling Layer: POOL

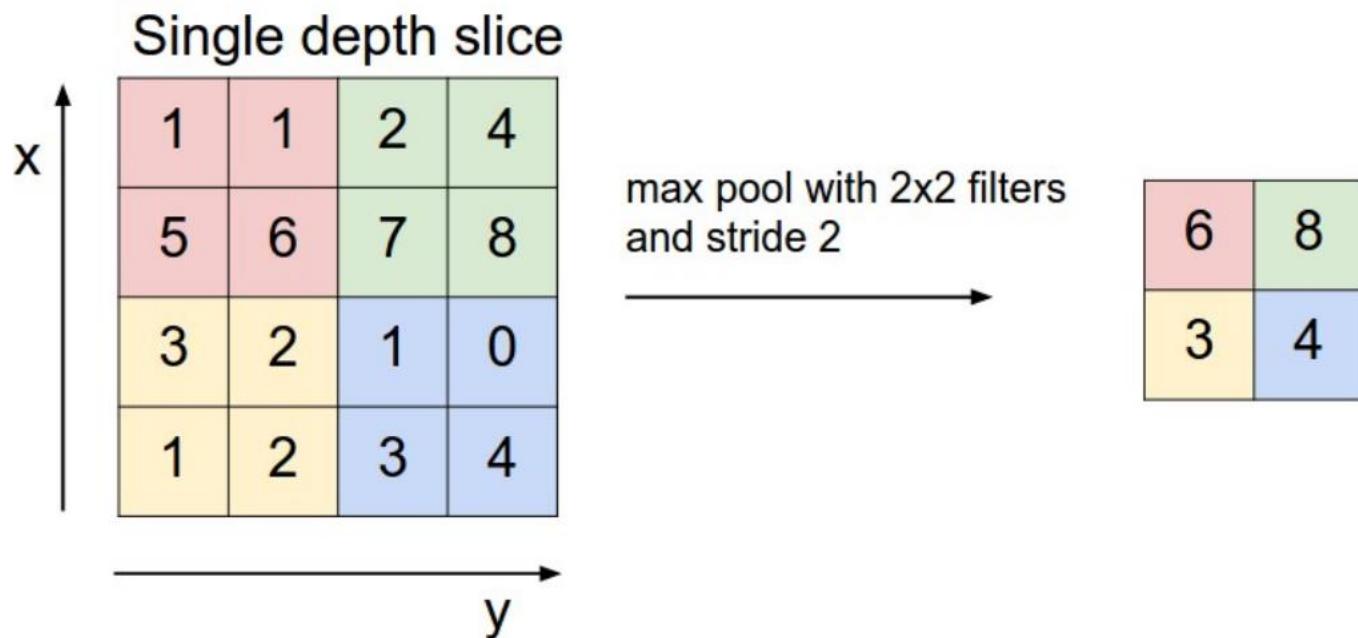
- Vastly more simpler than CONV
- Reduce the **spatial** size by using a MAX or similar operation
- Operate independently for each depth slice

POOL: Example

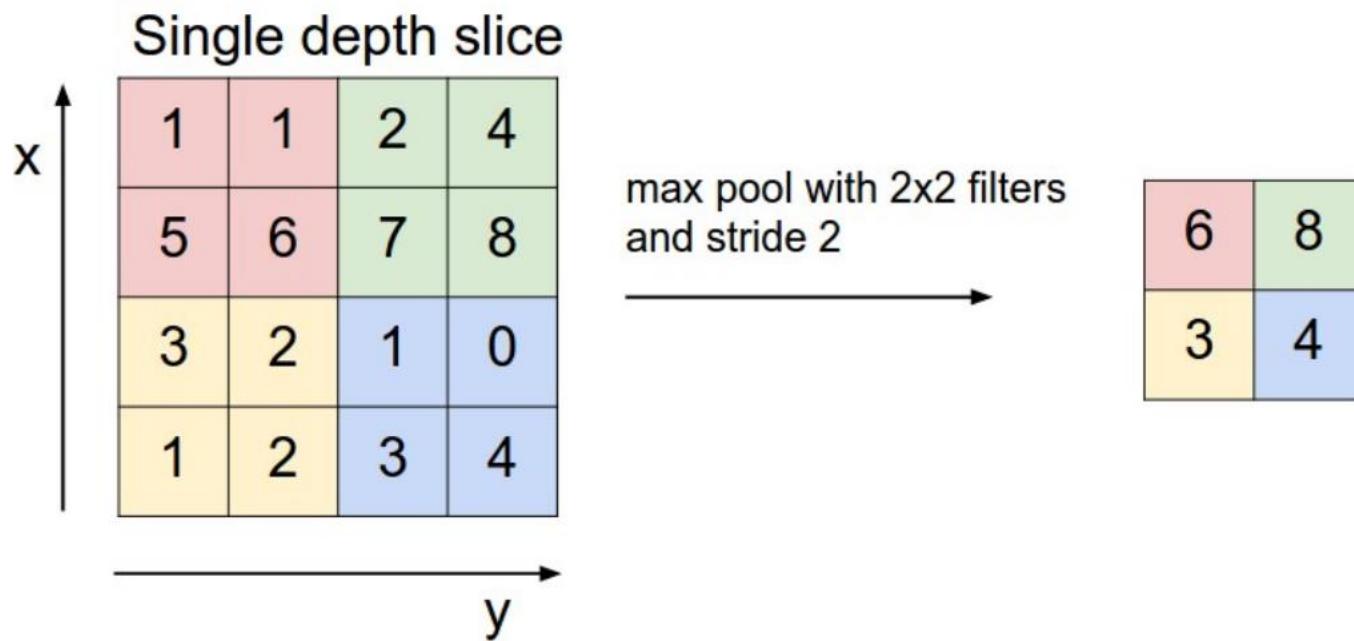
- Input depth is retained



POOL: Example



POOL: Example



- Recent research is showing that you may not need a pooling layer

Fully Connected Layer: FC

- Essentially a fully connected layer
- Already seen while discussing feedforward neural networks

Summary

- Feedforward neural nets can do better than linear classifiers (saw this for a low-dimensional small synthetic example)
- CNN have been very effective in image related applications.
- Exploit specific properties of images
 - Hierarchy of features
 - Locality
 - Spatial invariance
- Lots of **design choices** that have been empirically validated and are intuitive. Still, there is room for improvement.

Appendix

Naming: Why ‘Convolution’

