

---

---

# Deep Learning and Applications

Theja Tulabandhula

# Today's Outline

---

- Visualizing CNNs
- Transfer Learning
- Neural Net Training Tricks
  - Data Augmentation
  - Weight Initialization
  - Batch Normalization/Dropout

---

---

# Remarks: Convolutional Neural Networks

# CNN Architecture

- Typically a CONV is followed by a POOL
- Closer to the output, use FC layers
- In CONV, smaller filters are preferred (say  $3 * 3 * Z$ )
- Input image should ideally be divisible by 2 many times



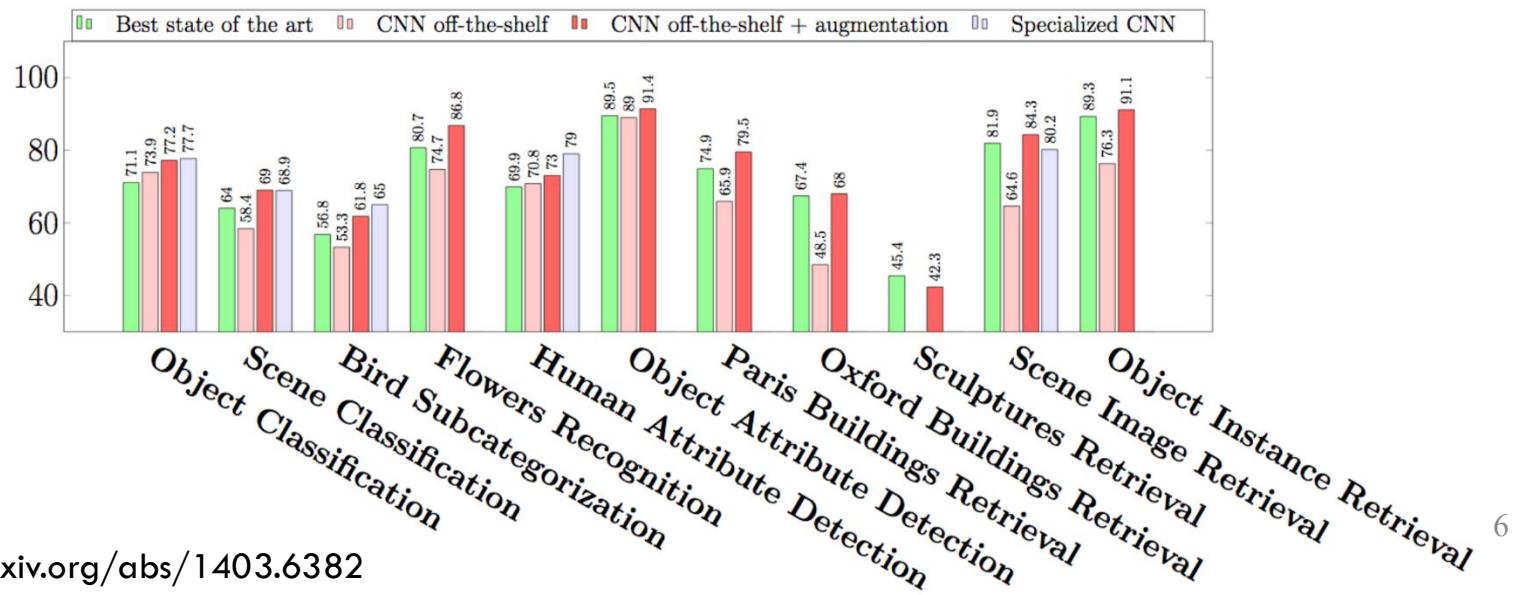
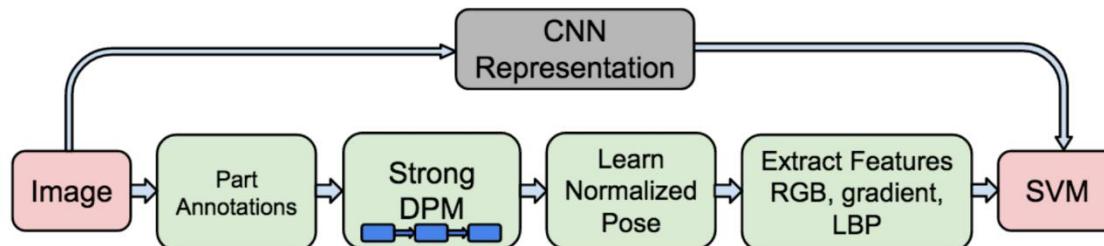
# VGG Net Example

- 2<sup>nd</sup> in the 2014 ILSVRC classification task
- 3x3 conv filters with stride 1
- ReLU non-linearity
- 5 POOL layers
- 3 FC layers



# Representational Power

- We can get a significant boost in performance compared to hand engineered classification/machine-learning pipelines



<sup>1</sup>Figure: <https://arxiv.org/abs/1403.6382>

# Example: CONV Layer Parameter Count

---

- Input tensor of size  $90 * 90 * 10$
  - Say we have 5 filters, each is  $3 * 3 * 10$
  - Stride is 1 and zero padding is 1
  - Then output tensor will be  $90 * 90 * 5$
  - We can calculate manually for other strides and padding values
- 
- Number of parameters is  $5 * (3 * 3 * 10 + 1) = 455$
  - Contrast with Fully connected net:
    - Number of inputs is 81000
    - Number of hidden layer neurons is 40500
    - Hence, the number of parameters is  $> 3,280,500,000$

# CNN and Backpropagation

---

- Backpropagation through a CONV layer
  - Constitutes a set of matrix-matrix products and whatever is the behavior for the nonlinearity
- Backpropagation through a POOL layer
  - Essentially like ReLU where one can keep track of the index of the maximum

---

---

# Questions?

# Visualizing CNNs

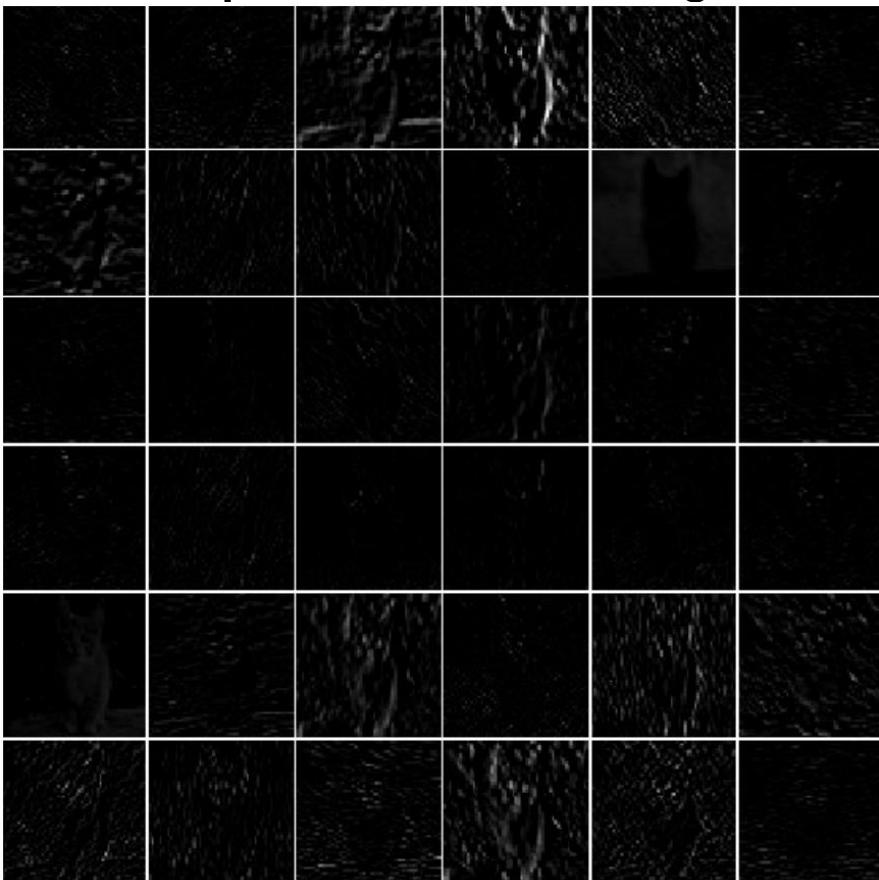
# Combating Non-Interpretability

---

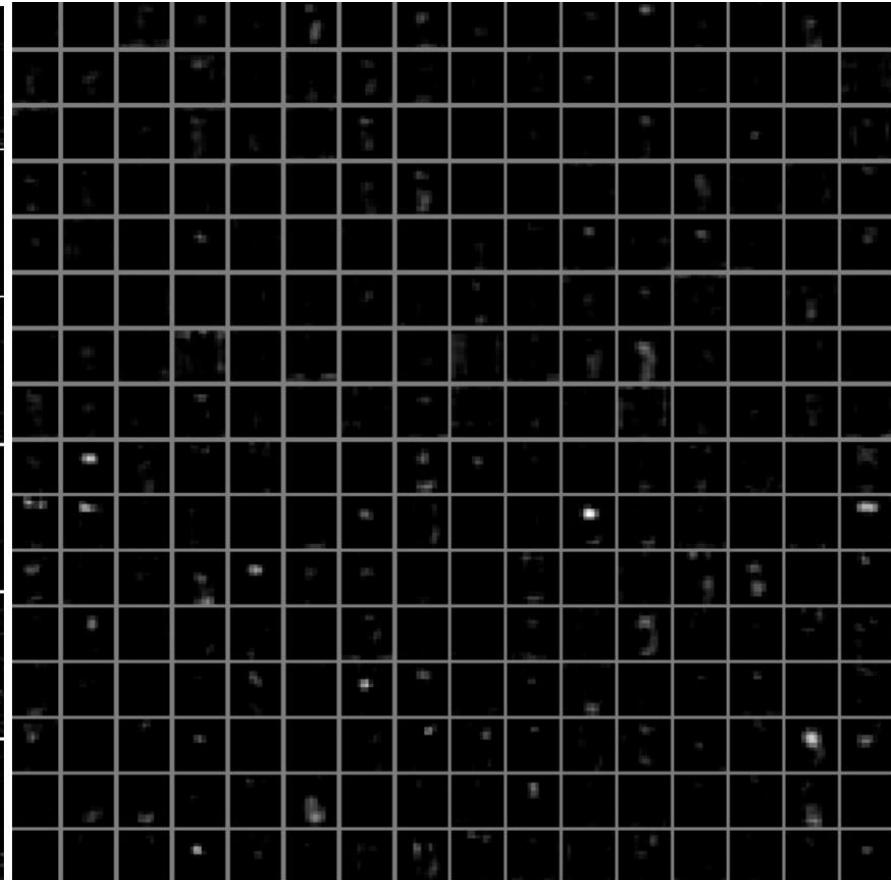
- Common criticism: learned features are not interpretable
- We will make a few attempts
  - Look at activations
  - Look at weights
  - Look at images in an embedded space
  - Look at impact of occlusion
  - Look at images that activate neurons highly

# Visualize: Activations

- Useful to debug ‘dead’ filters (e.g., when using ReLU)
- Input is a cat image



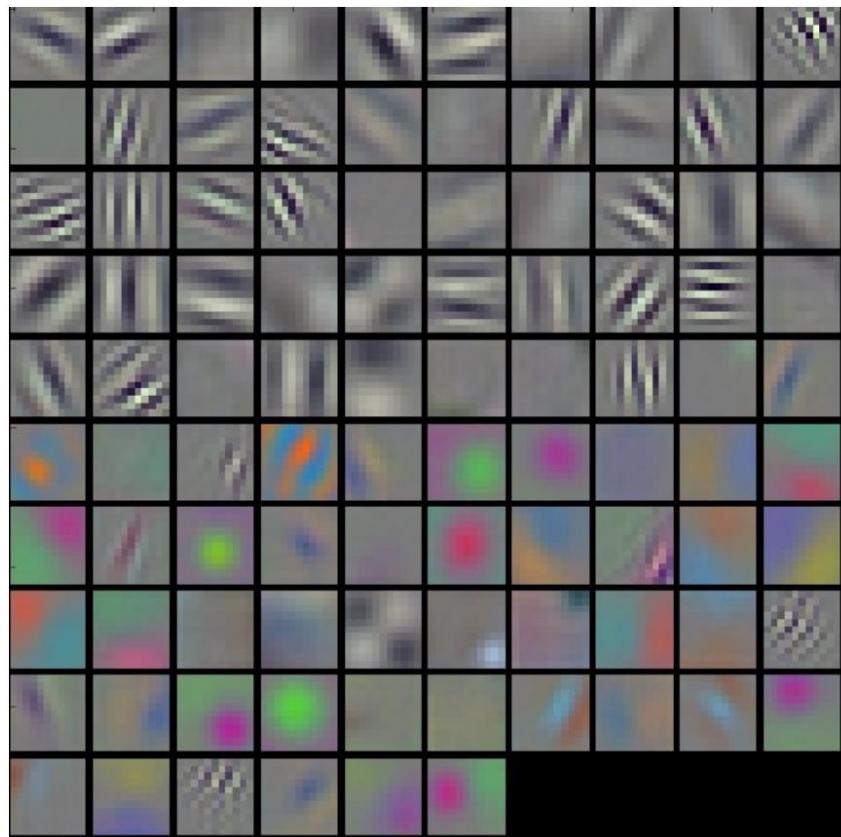
1<sup>st</sup> CONV



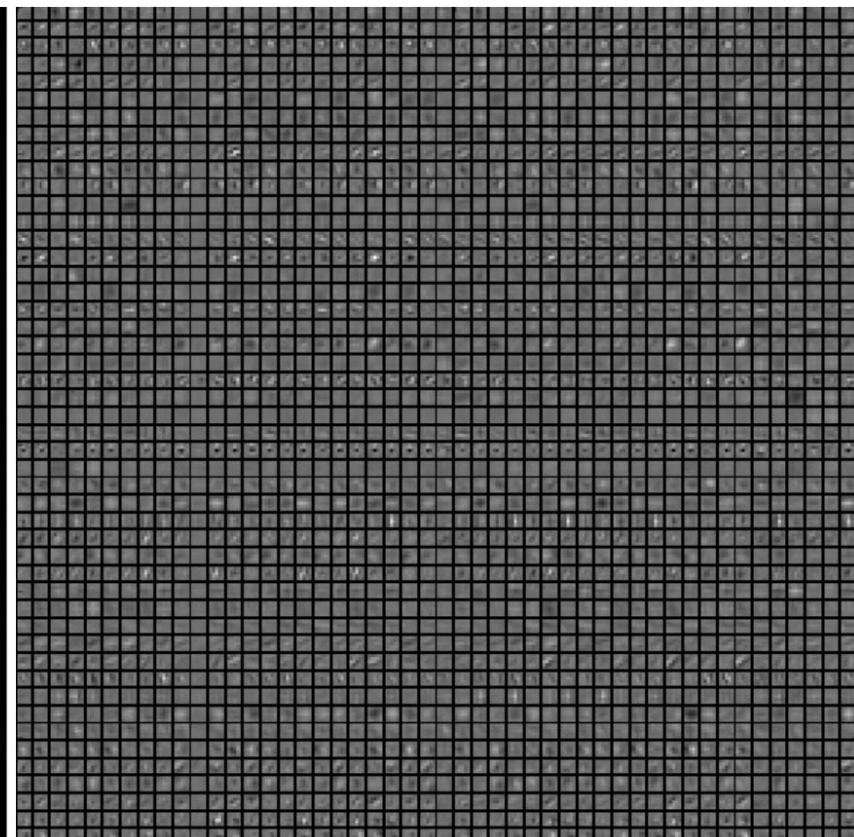
5<sup>th</sup> CONV

# Visualize: Weights

- Useful to debug if training needs to be run more (if patterns are noisy)



1<sup>st</sup> CONV



2<sup>nd</sup> CONV

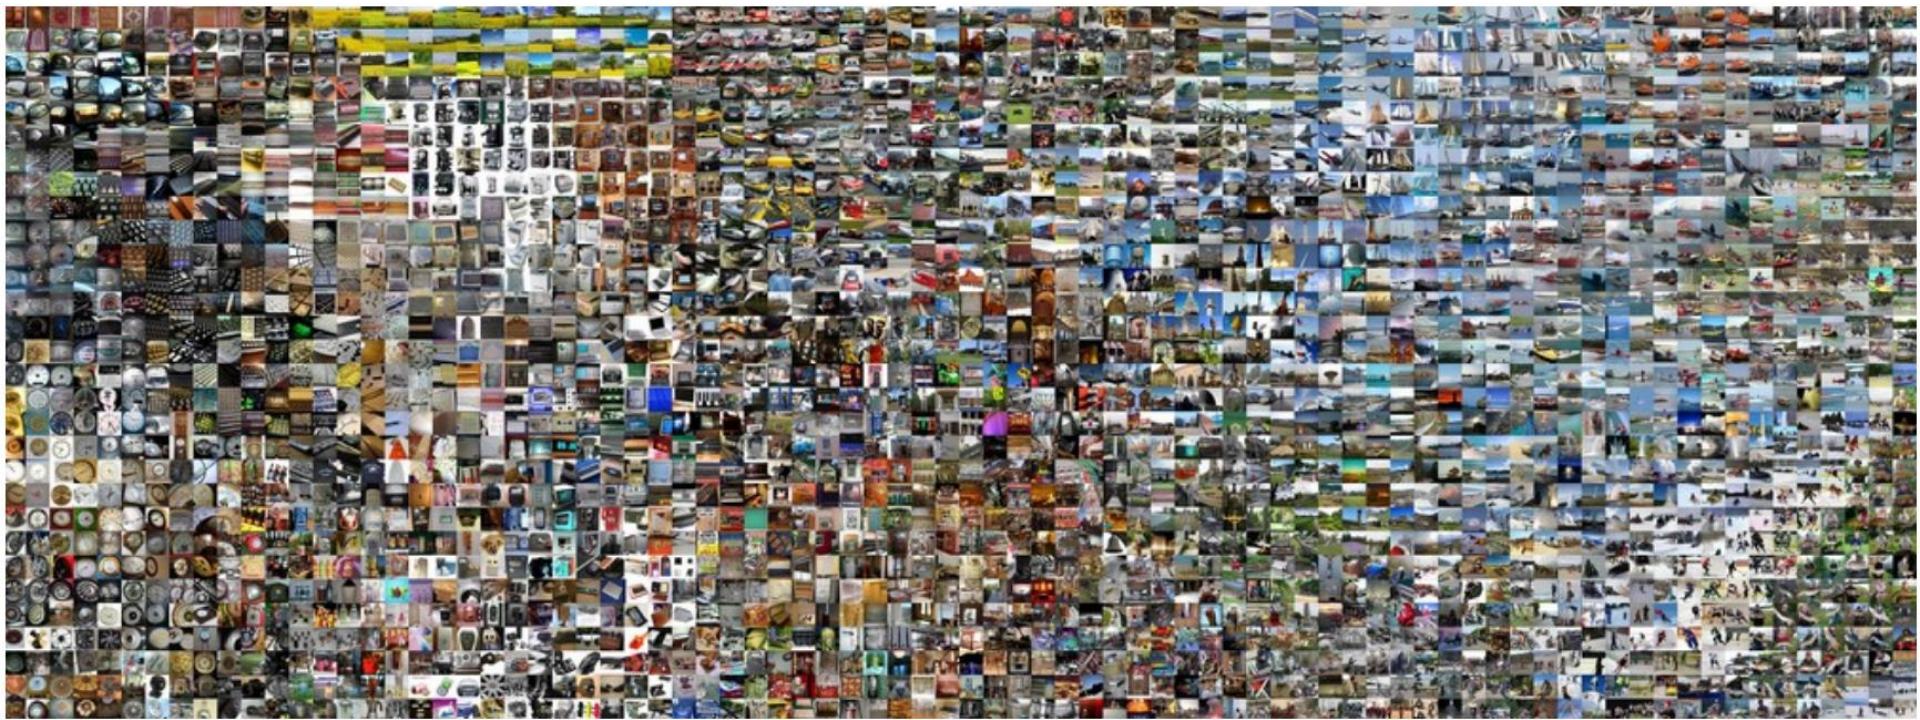
# Visualize: Low-Dimensional Embeddings

---

- CNN
  - Input: Image
  - Output: Scores
- The input to the layer that computes scores:
  - $s = W \max(0, h) + b = Wa + b$
- Activation  $a$  can be considered as a representation of the input image
- Embed  $a$ 's into a 2D space
  - Such that distance properties are preserved

# Visualize: Low-Dimensional Embeddings

- In Alexnet, the output of layer before FC layer is 4096 dim
- The t-SNE embedding is shown below:



- Similarities are class-based and semantic rather than color and pixel based
  - Implies: images close to each other are similar for the CNN<sub>15</sub>

<sup>1</sup>Figure: <http://cs231n.github.io/understanding-cnn/>

# Visualize: By Occlusion

---

- To figure out which part of the image is leading to a certain classification
- Plot the probability of class of interest as a function of occlusion

# Visualize: By Occlusion

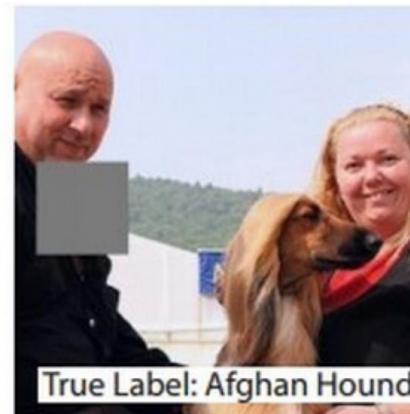
- Occlusion in grey is slid over the images and plot probability of correct class



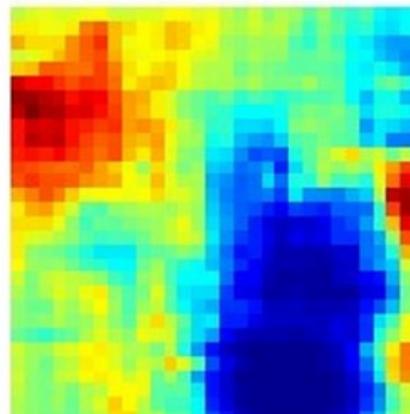
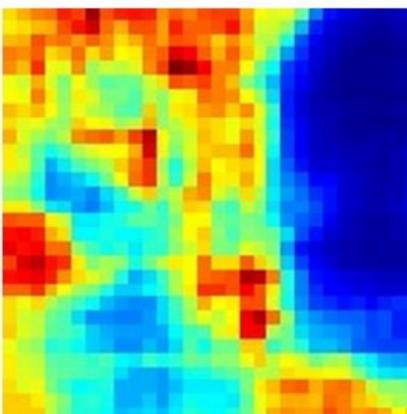
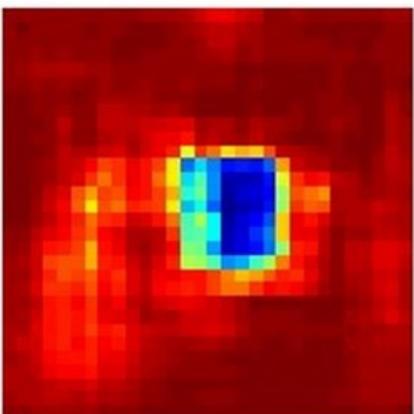
True Label: Pomeranian



True Label: Car Wheel

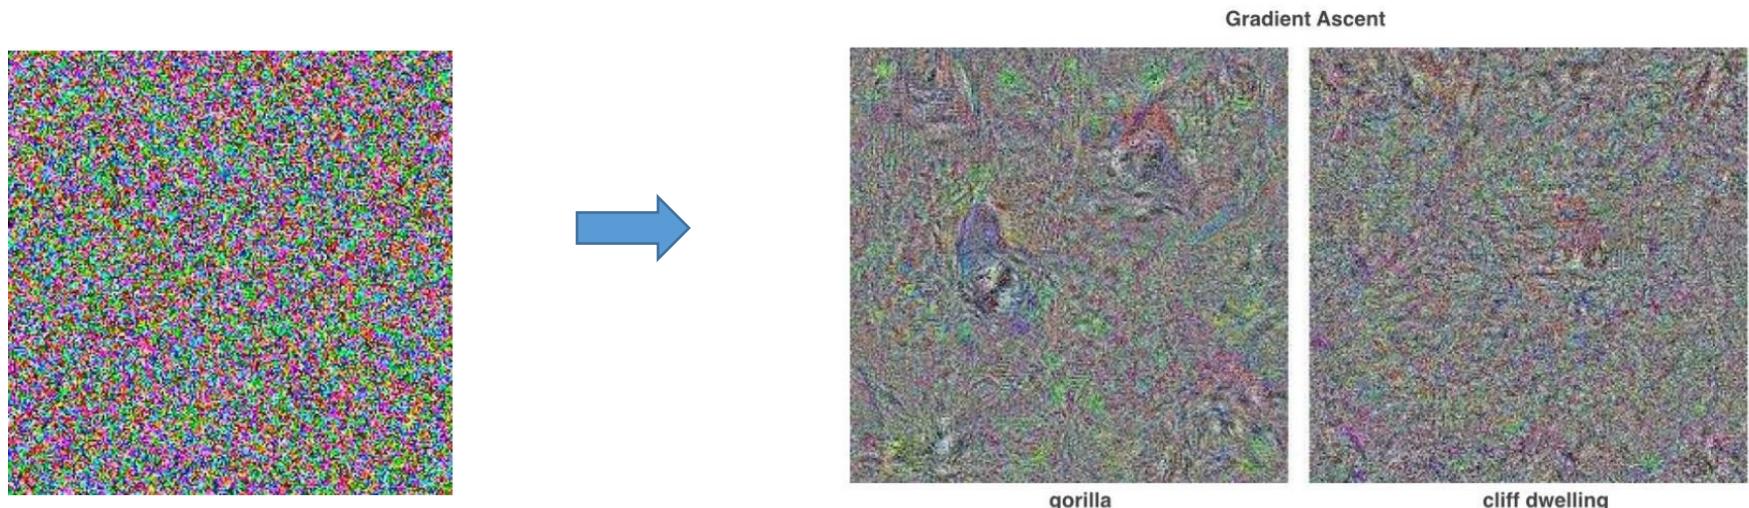


True Label: Afghan Hound



# Visualize: Synthesize Images

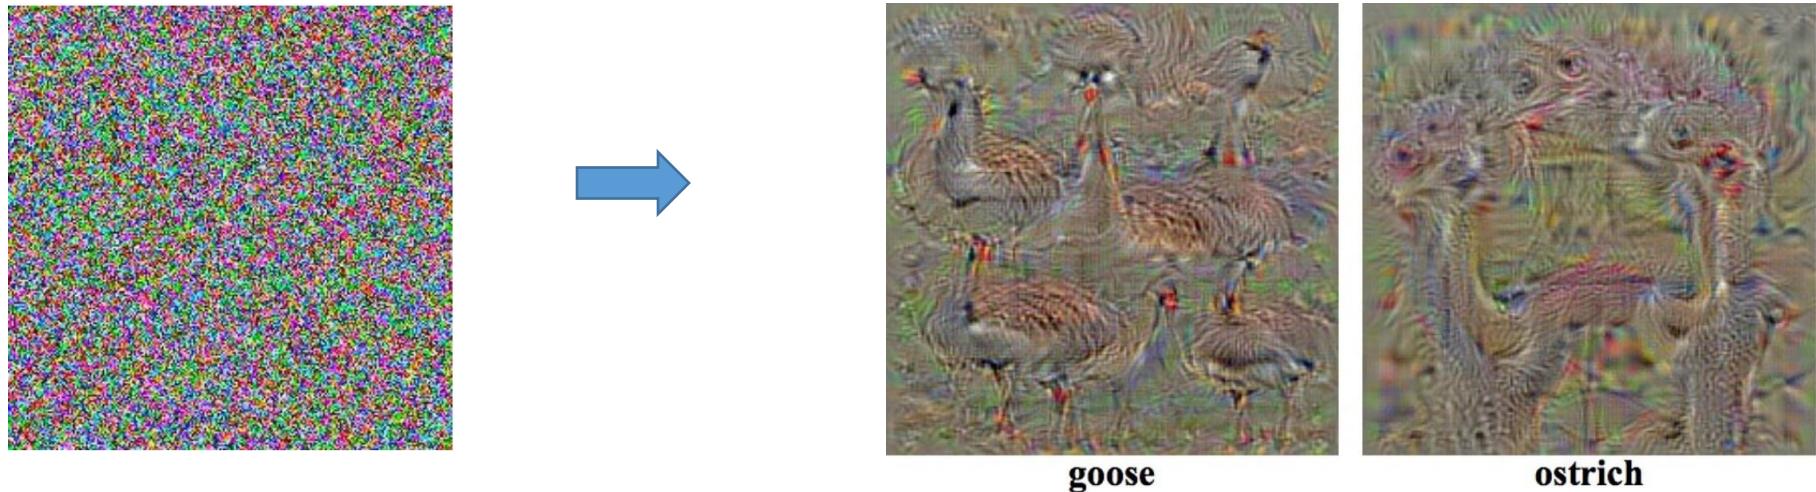
- Find images that activate a neuron the most



- Seed with ‘natural’ image priors

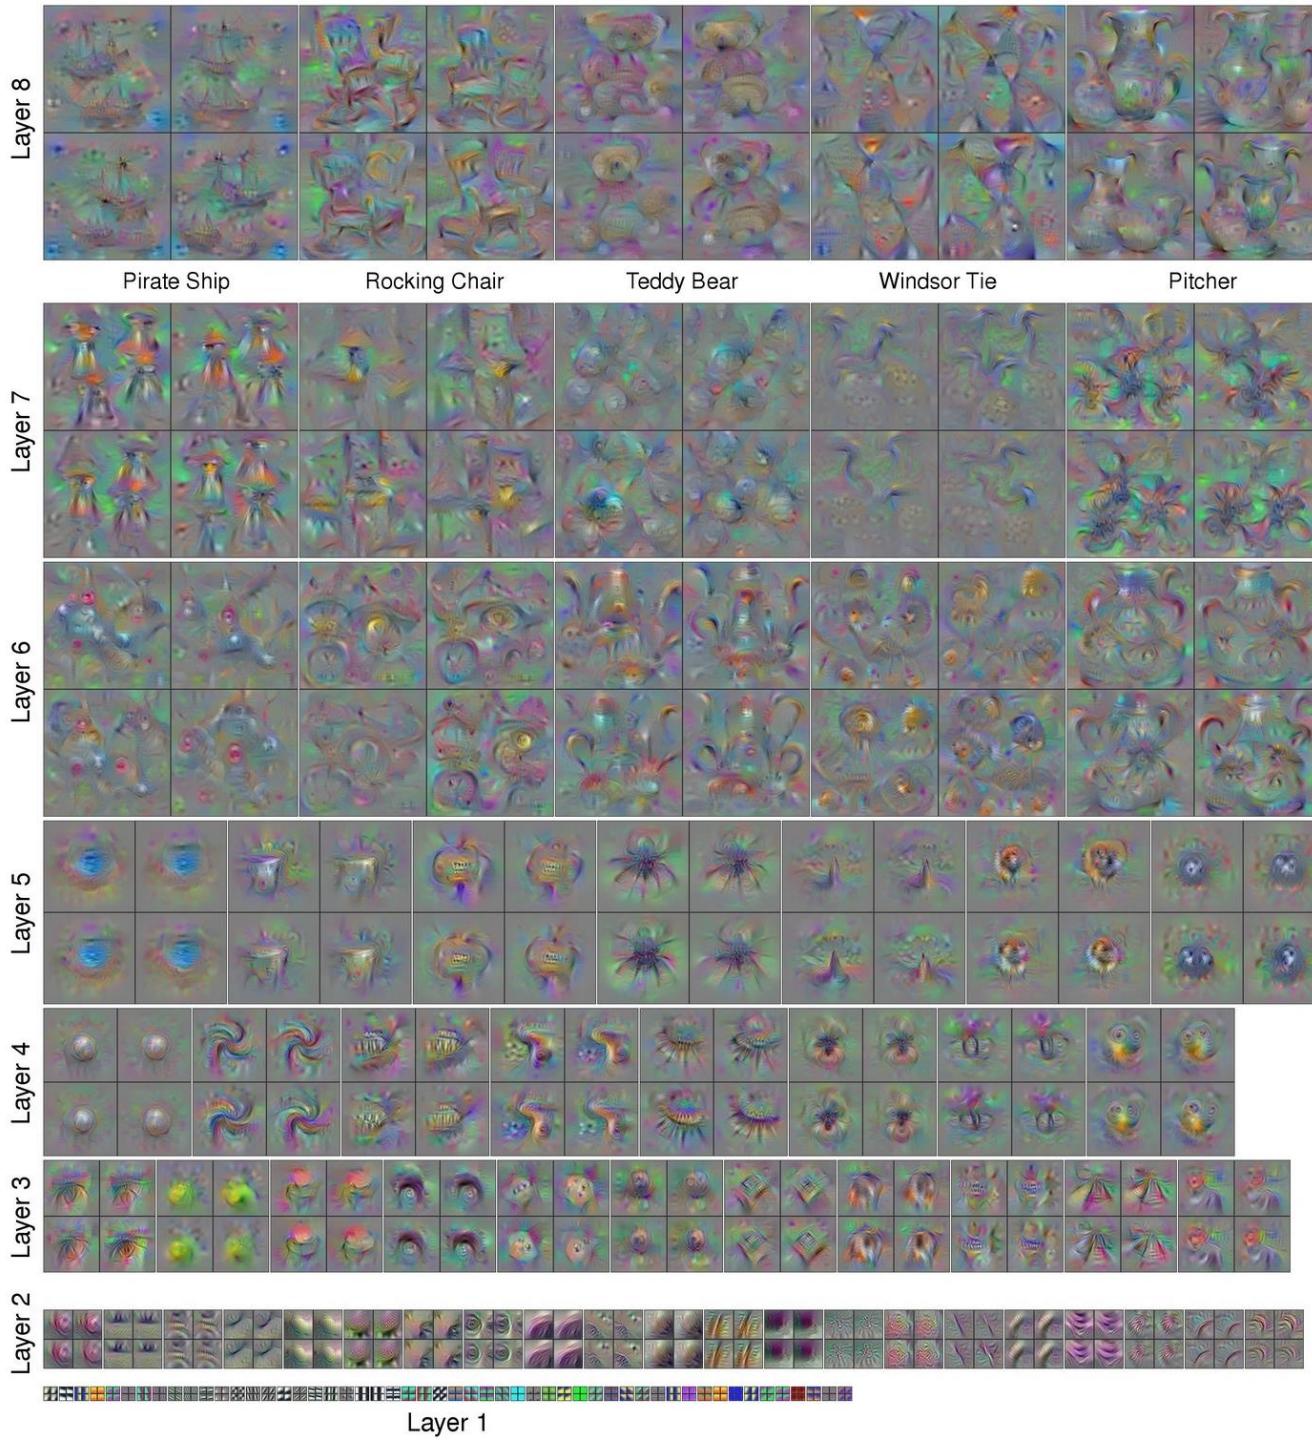
# Visualize: Synthesize Images

- Find images that activate a neuron the most



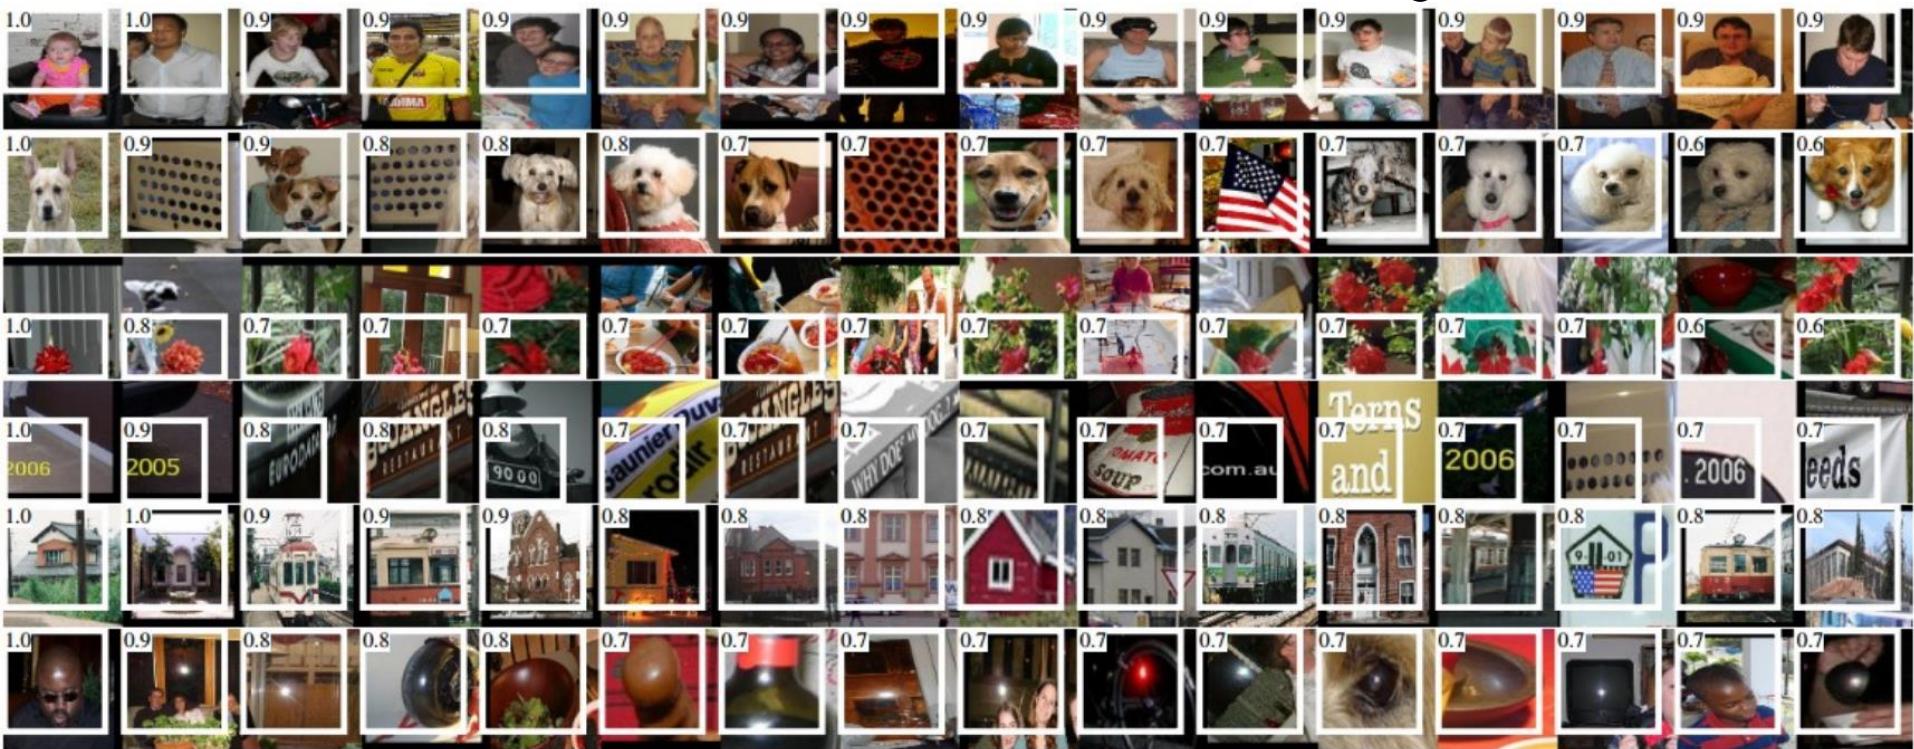
- Seed with ‘natural’ image priors

# Visualize: Synthesize images



# Visualize: Images that Activate a Neuron

- Track which images maximally activate a neuron
  - Understand what the neuron is tracking



5<sup>th</sup> POOL      Activation values and receptive fields of some neurons in Alexnet  
(May not be a good idea...)

---

---

# Questions?

# Today's Outline

---

- Visualizing CNNs
- Transfer Learning
- Neural Net Training Tricks
  - Data Augmentation
  - Weight Initialization
  - Batch Normalization/Dropout

---

---

# Transfer Learning

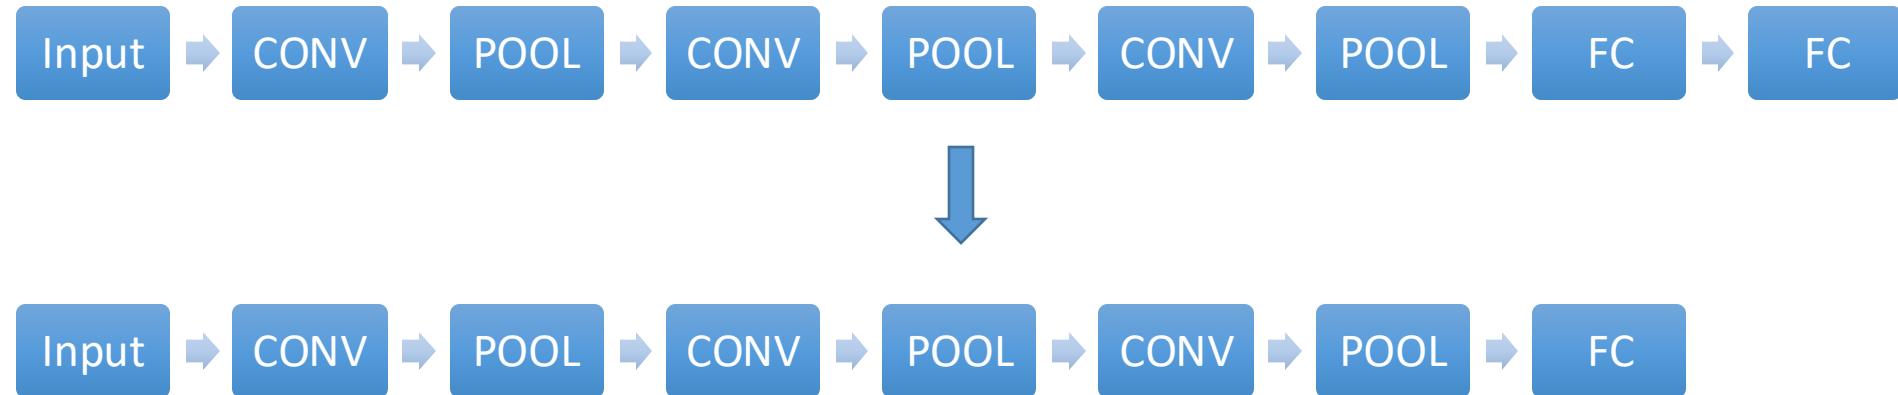
# Transfer Learning

---

- Very few people train a deep feedforward net or a CNN from scratch
- **Myth:** “We need a lot of data to use Deep Neural Networks”
- We will see two approaches if we have small data
  - Feature extraction
  - Fine-tuning (including LoRA fine-tuning)
- Both these are loosely termed as **Transfer learning**

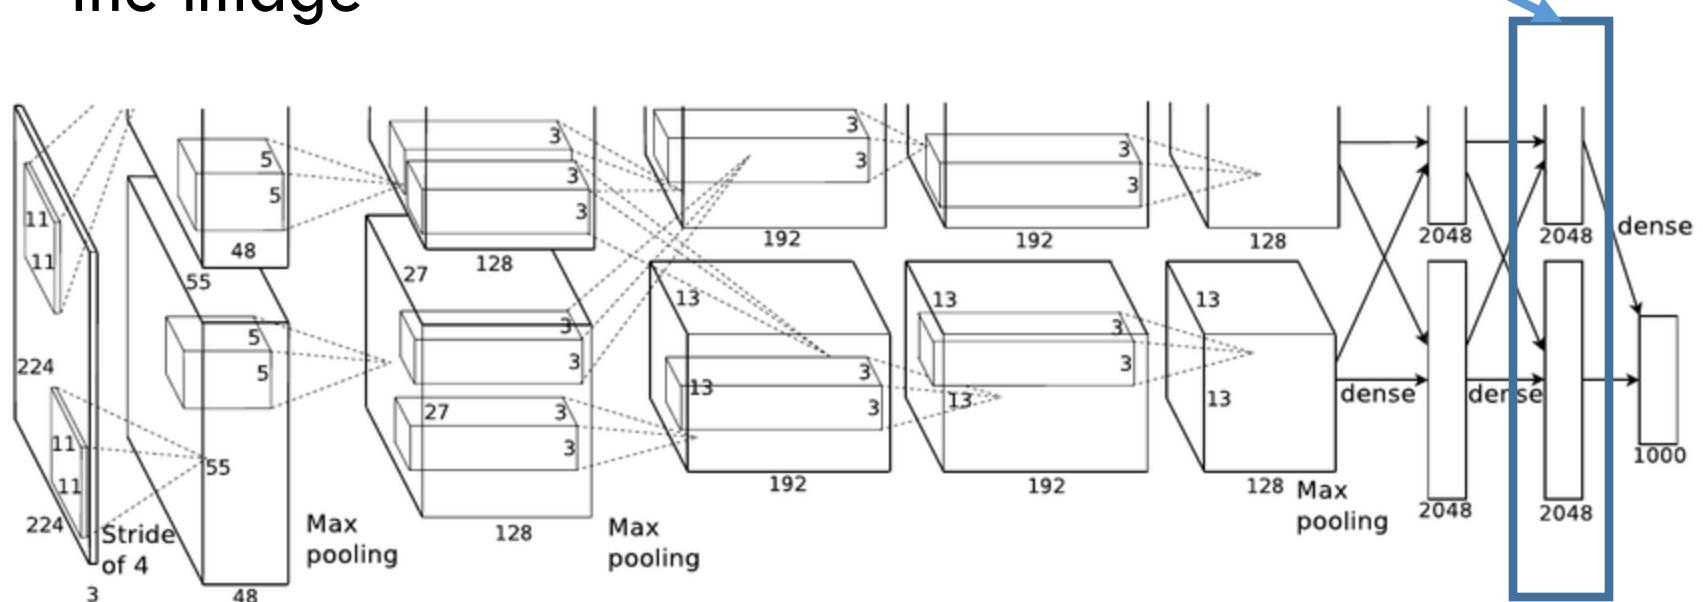
# Transfer by Feature Extraction (I)

- Get a pretrained CNN
  - Example: VGG or AlexNet that was trained on Imagenet
- Remove the last FC (that outputs 1000 dim score)
- Pass new training data to get **embeddings**



# Image Embeddings

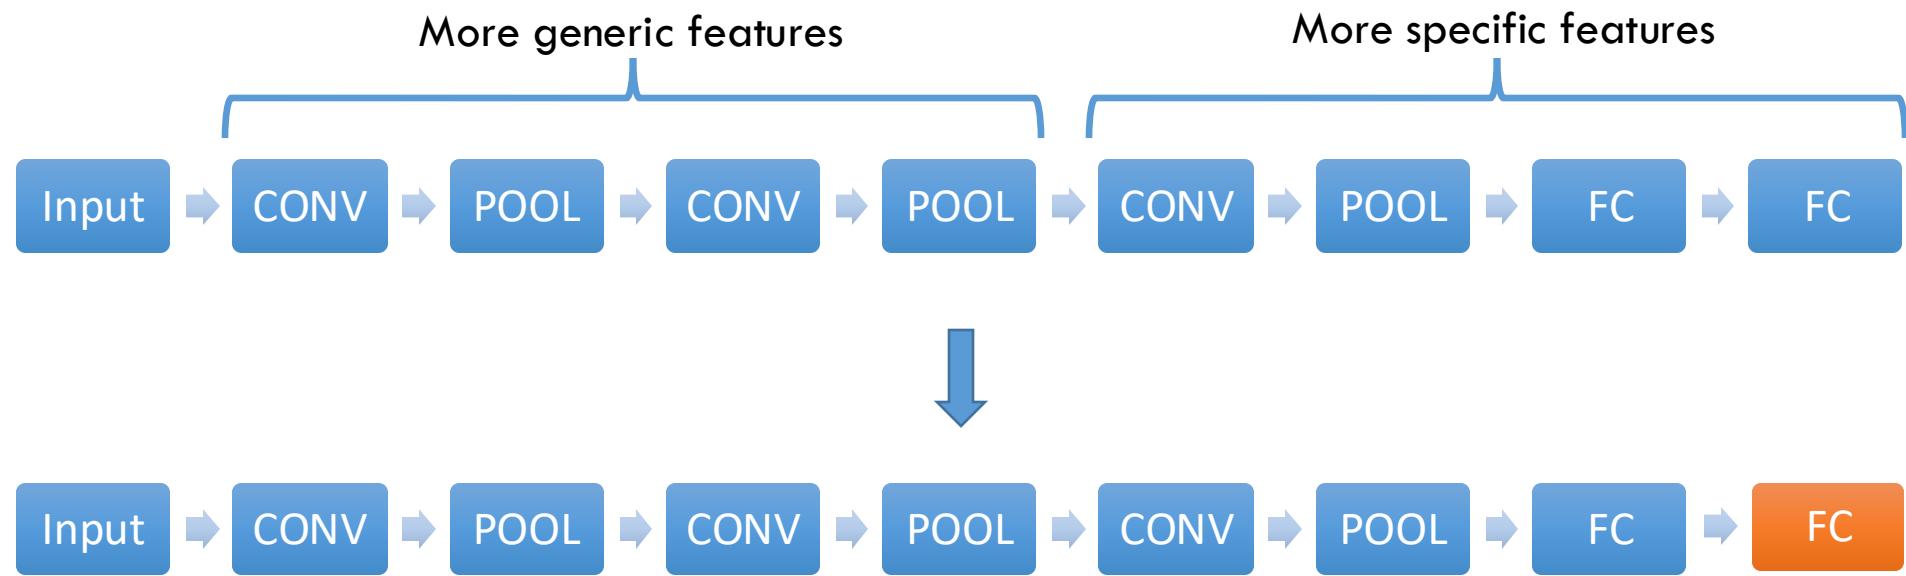
- We can think of the penultimate hidden layer activations (a 4096 dim vector) as an **embedding** of the image



- This is the **activation vector** or the **representation** or the **CNN code** of the image

# Transfer by Feature Extraction (II)

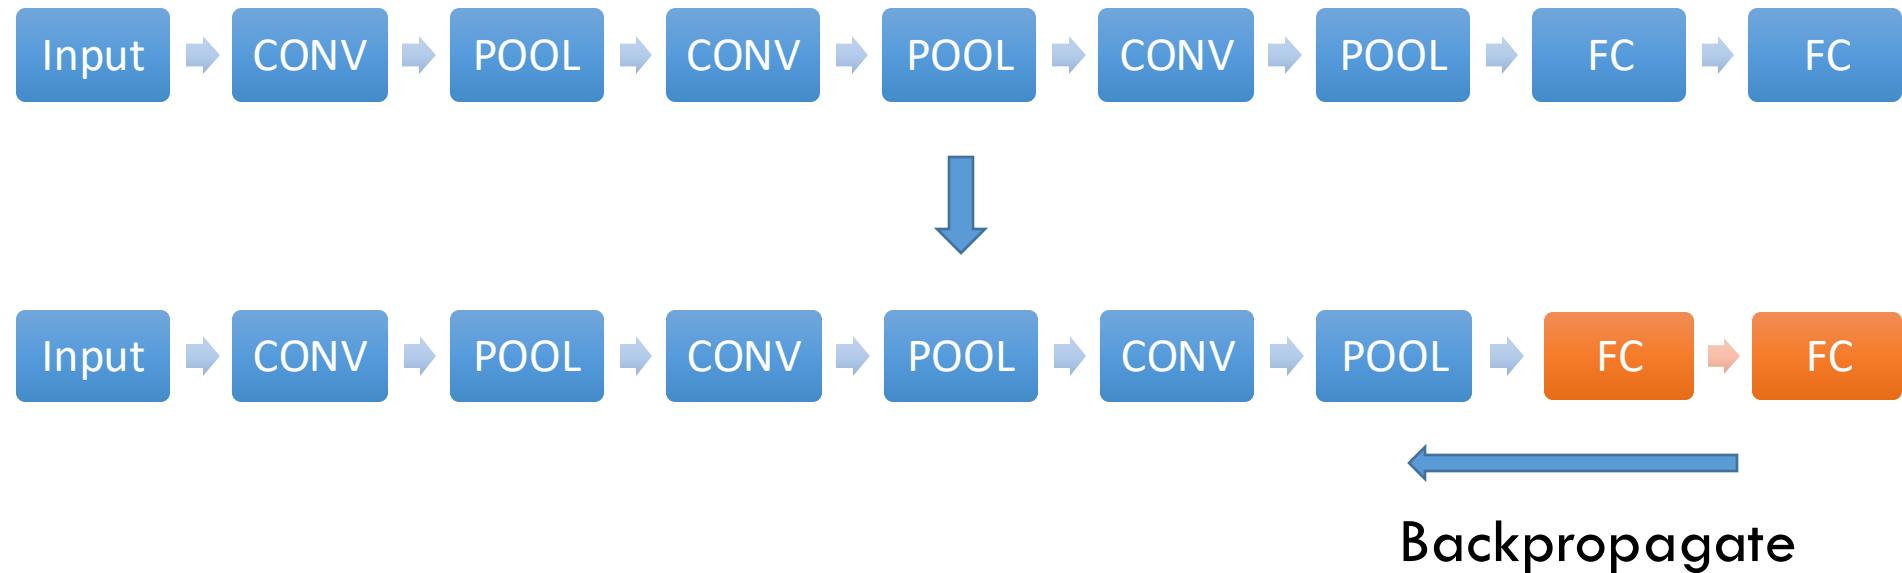
- Input these to a linear or non-linear classifier!



- For example, for imagenet output 1000 dim scores
- For our data, output say 2 scores (cat vs dog)

# Transfer by Fine-tuning

- Retrain or **finetune** additional layers of the pre-trained if we have more data



- We can even go all the way back to the first layer if there is a lot of training data available

# Transfer by Parameter Efficient Fine-tuning

---

- Finetuning can be made parameter efficient
- One popular way to do this is called LoRA (2021)
- Enables a lot of organizations to adapt very capable general-purpose models to specific applications and achieve non-trivial gains

# Transfer Learning Choices

- When to transfer

	Similar dataset	Different dataset
Small data	Feature extract	NA
Large data	Fine-tune a bit	Fine-tune a lot

- How to transfer
  - Get pre-trained models for popular software systems (e.g., from Hugging Face Hub)

This is key for projects!

# Aside: Other Vision Tasks

- Some example vision tasks are given below

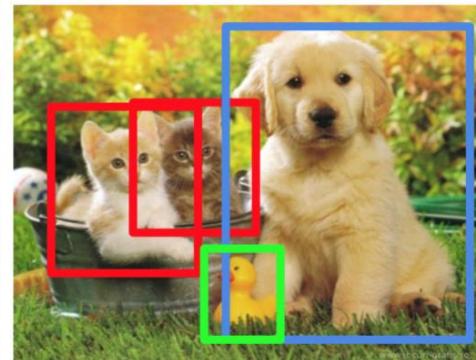
**Classification**



**Classification + Localization**



**Object Detection**



**Instance Segmentation**



CAT

CAT

CAT, DOG, DUCK

CAT, DOG, DUCK

Single object

Multiple objects

---

---

# Questions?

# Today's Outline

---

- Visualizing CNNs
- Transfer Learning
- Neural Net Training Tricks
  - Data Augmentation
  - Weight Initialization
  - Batch Normalization/Dropout

---

---

# Neural Net Training Tricks

# Neural Nets in Practice

---

- There are a few **empirically validated** techniques that improve the performance (classification accuracy) of feedforward nets and CNNs
- We will look at some of these
  - Data: data augmentation
  - Model: initialization, batch normalization, dropout
- For our discussion, we will fix the optimization technique to be a gradient based method. We will revisit related algorithmic enhancements later.

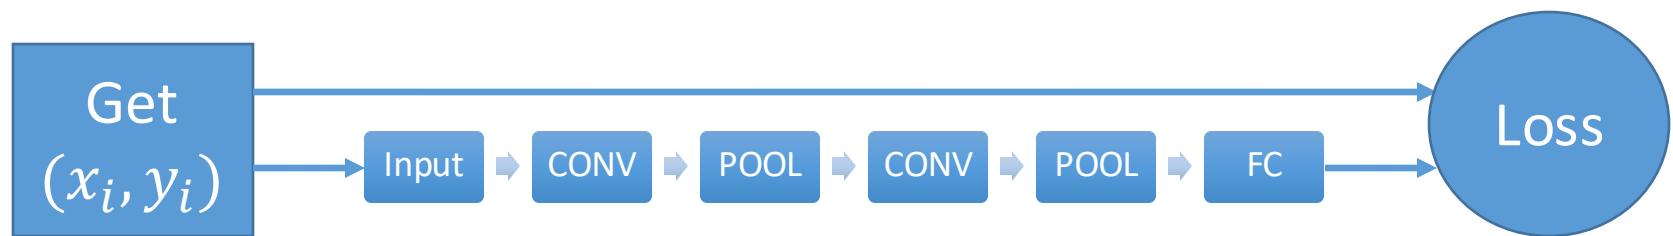
# Data

---

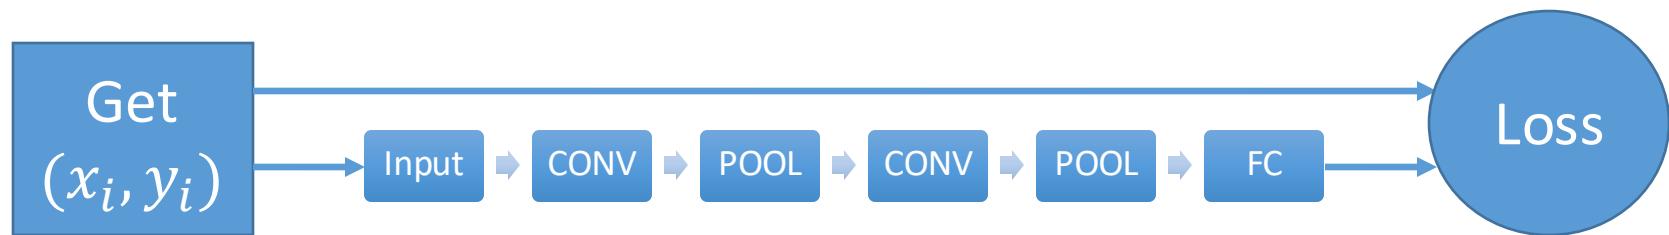
- Data:
  - How is it handled?
  - What is its quality?
- Handling:
  - Deep nets may need to read lots of data (images), so keep them in contiguous spaces of hard-disk
- Quality:
  - Collect as much clean data as possible. At the same time, unclean may also be good enough

Next: Extract the most out of existing data for CNNs

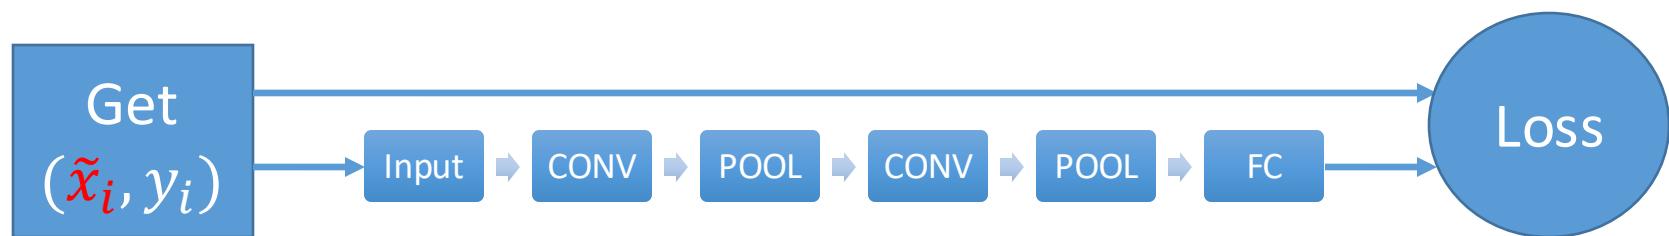
# Augmenting Data (I)



# Augmenting Data (I)



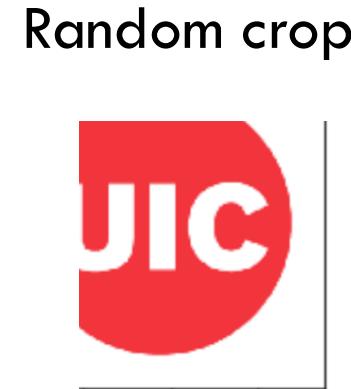
And



Where  $\tilde{x}_i = g(x)$  is a transformation

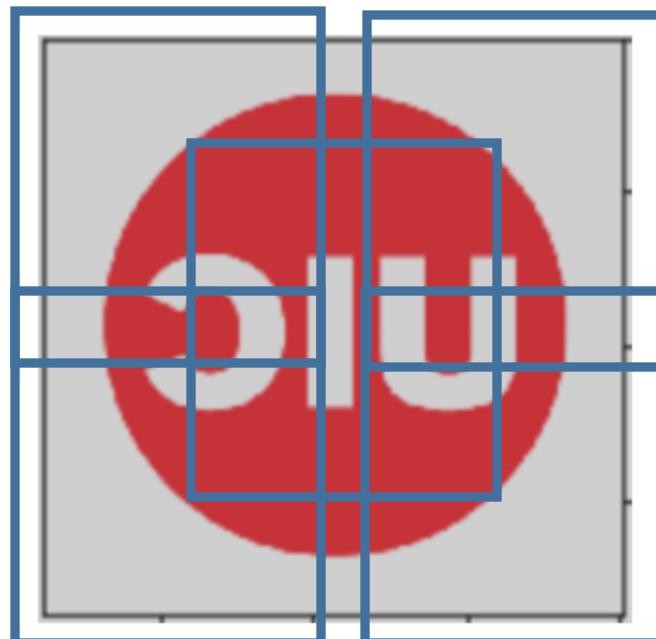
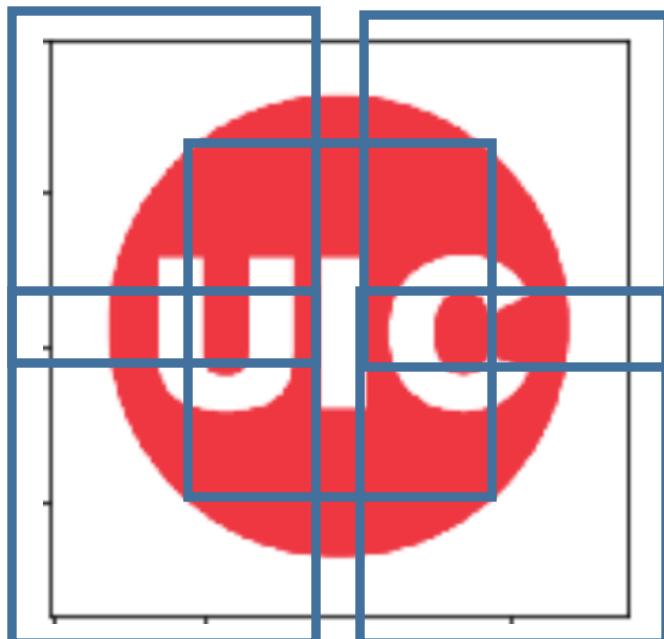
# Augmenting Data (II)

- We are changing the input without changing the label
- We then add this new example to our training set
- Widely used technique!



# Augmenting Data (III)

- At test time, **average** the predictions of a fixed set of transformations
- Example (for Resnet, the ILSVRC 2015 winner):
  - Image at 5 scales: 224,256,384,460 and 640
  - At each scale, get 10 224\*224 crops



# Augmenting Data (IV)

---

- Other ways to augment data include
  - Changing contrast and color
  - Mix translations, rotations, stretching, shearing, distortions
- This is very useful for small datasets
- From one point of view, this is essentially
  - Adding some noise during training
  - Marginalizing noise out at test

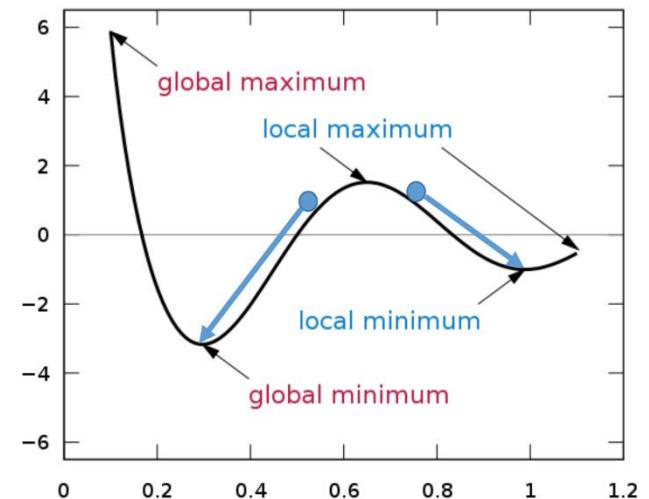
# Optimization

---

- We have already seen few choices
  - Learning rates
  - Optimization schemes: stochastic gradient descent (SGD), SGD with momentum, Adam, RMSProp
- We briefly discuss on other choice while training deep neural nets (including CNNs) that can make a difference
  - Weight initialization

# Optimization: Weight Initialization

- Weight initialization plays a key role in training deep networks
  - Example:  $W = 0$  may be bad
- Not just the issue of local optima
- But also the magnitudes of gradients in backprop
  - Activation statistics (mean and variance) influence gradients
- Heuristics available in the literature to initialize  $W$



# Model

---

- We have already seen few choices
  - Activation function or nonlinearities
  - Number of layers and number of neurons per layer
  - CNN filter choices ...
- There are other choices while training deep neural nets (including CNNs) that also **make a difference**
  - Batch normalization
  - Dropout

# Model: Batch Normalization

---

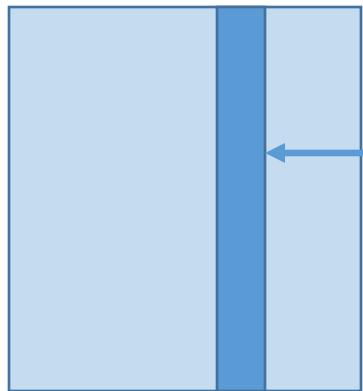
- Activations magnitudes and their statistics depend on the dataset, the network and the nonlinearity used
- Their statistics influence gradient propagation, hence also learning
- Is there a way to control them?
  - Yes, through batch normalization!

# Model: Batch Normalization

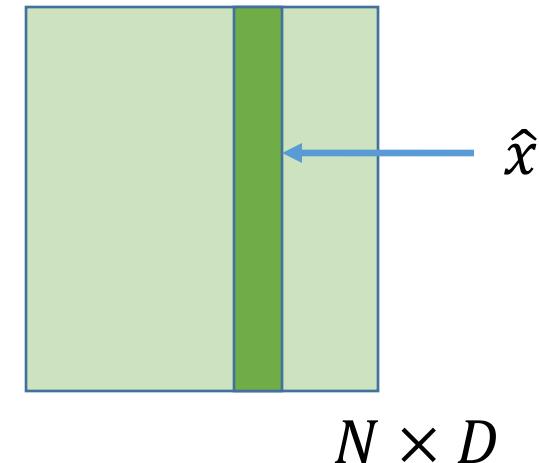
- Idea: Make each activation unit-Gaussian by subtracting the mean and then dividing by standard deviation

Batch-size =  $N$

Number of output neurons =  $D$



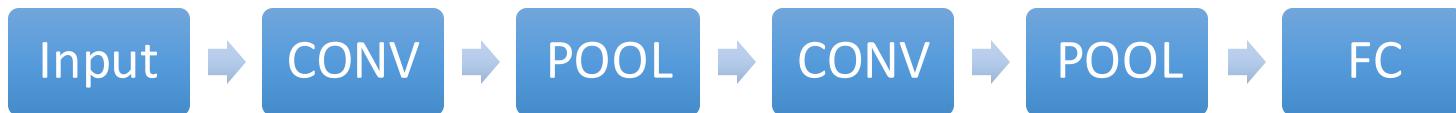
$$\hat{x} = \frac{\gamma(x - E[x])}{\sqrt{Var[x]}} + \beta$$



- Is a differentiable function: hence no issue with backpropagation
- At test time, there is no batch. Use the training data means and variances

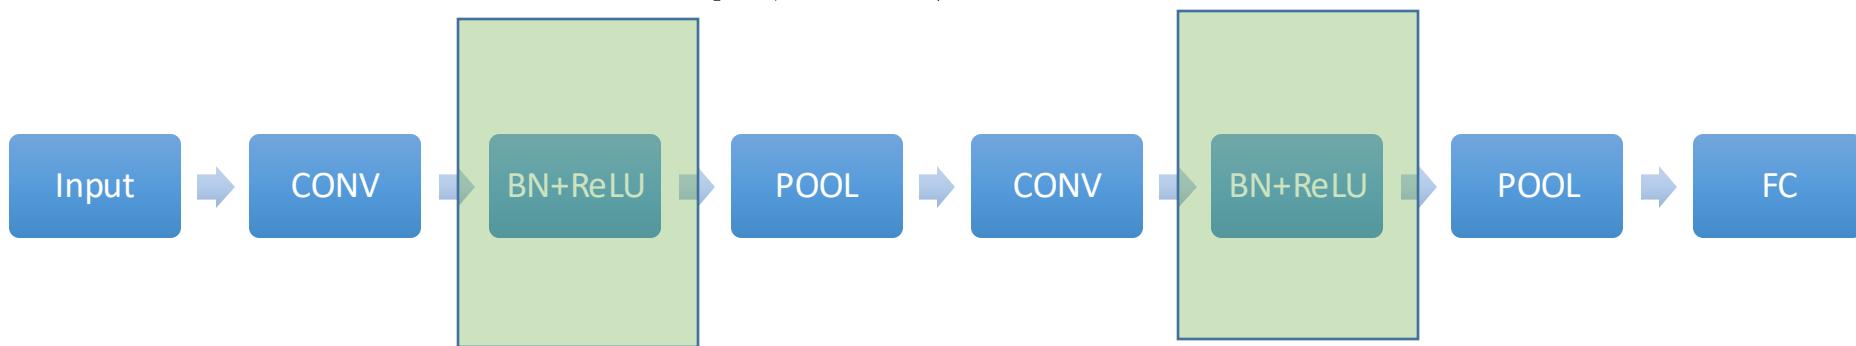
# Model: Batch Normalization

- Previously,



- Now

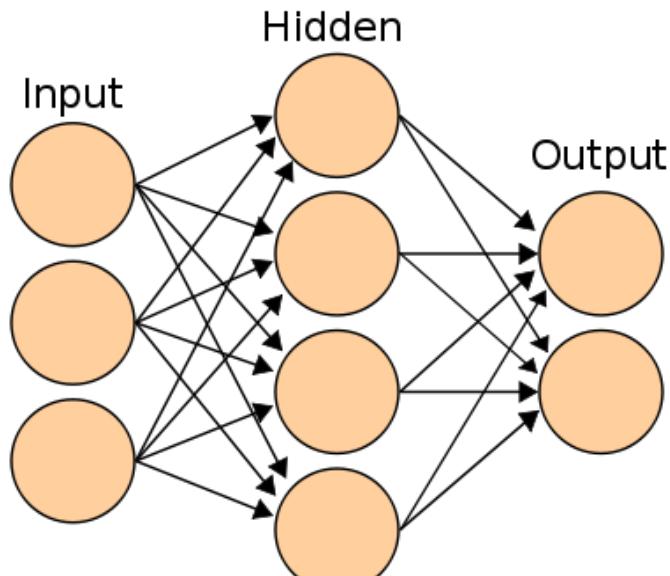
- Insert a Batch Normalization layer between CONV and nonlinearity (ReLU)



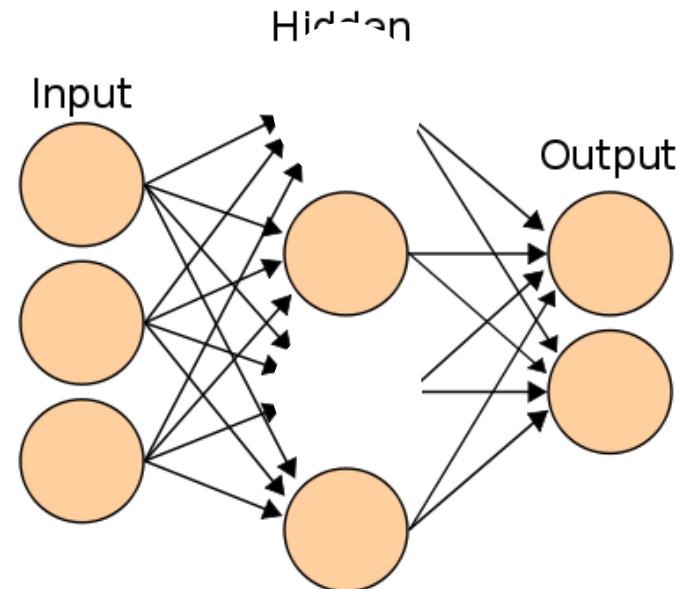
- Empirically observed: improved gradient flows, less sensitive to **initialization**.

# Model: Dropout (Regularization)

- Idea: During training, every time we forward pass, we set the output of a few neurons to zero with some probability



Without dropout



One pass with dropout

# Model: Dropout (Regularization)

---

- Intuitively, it is
  - Making us use smaller capacity of the network. Hence, can think of it as a regularization
  - Forcing all the neurons to be useful. Hence there is over-representation or redundancy
- Also think of it as
  - Subsampling a part of the network for each example
  - Thus, we get an ensemble of neural networks that share parameters

# Model: Dropout (Regularization)

---

- Higher probability means stronger regularization
- At test time,
  - Instead of doing many forward passes
  - Perform no dropout
  - Scale all activations by the probability of dropout
- Example:
  - Say dropout with probability  $p$
  - Originally:  $f(x, W_1, b_1, W_2, b_2) = W_2 \max(0, W_1 x + b_1) + b_2$
  - With dropout:  $W_2 * p * \max(0, W_1 x + b_1) + b_2$

# Summary (I)

---

- CNN are very effective in image related applications.
  - State of the art!
- Exploit specific properties of images
  - Hierarchy of features
  - Locality
  - Spatial invariance
- Lots of **design choices** that have been empirically validated and are intuitive. Still, there is room for improvement.

# Summary (II)

---

- We saw
  - Visualizations to understand how CNNs work
  - Transfer learning applied to CNNs (important for applications)
    - An excellent way to get a deep learning solution working
    - There is no need for large datasets to get started

# Summary (III)

---

- Neural Nets Training Tricks
  - Revisited data: data augmentation
  - Revisited models: initialization, batch norm, dropout
- To train state of the art deep learning systems, you have to rethink:
  - (a) data, (b) models, **loss**, and (c) optimization<sup>1</sup>
  - What is the most bang per buck for your business?
- If the deep learning system is core to the business, look at engineering best practices (we saw some today)

<sup>1</sup>We did not cover loss function design in this lecture