

# **Introduction to PyTorch Programming Patterns**

Deep Learning and Modern Applications

---

Theja Tulabandhula

University of Illinois Chicago



After this lecture, you will be able to:

- Use PyTorch's device abstraction for GPU/CPU computation
- Build data pipelines using DataLoader and transforms
- Define models using the nn.Module class pattern
- Implement the standard training loop correctly
- Control gradient computation and model behavior modes
- Manage model state for saving, loading, and transfer learning

## The PyTorch Philosophy

- **Imperative programming:** Write code that executes immediately
- **Dynamic computation graphs:** Build graphs on-the-fly each forward pass
- **Pythonic:** Feels like native Python, uses standard debugging tools
- **Research-friendly:** Easy to experiment and modify

## PyTorch vs Other Frameworks

Aspect	PyTorch	TensorFlow 1.x	TensorFlow 2.x
Graph	Dynamic	Static	Eager by default
Debugging	Standard Python	<code>tf.Session</code> complexity	Improved
API Style	Object-oriented	Symbolic	Keras-like
Research Adoption	Very High	Moderate	Growing

PyTorch is now the **dominant framework** in research publications.

## The Device Abstraction

PyTorch uses a device abstraction to manage CPU/GPU computation:

```
import torch

# Check GPU availability
device = torch.device("cuda:0" if torch.cuda.is_available()
                      else "cpu")
```

This pattern appears in **every** PyTorch script.

## Moving Data to Device

```
# Create tensor on CPU
x = torch.randn(3, 4)

# Move to device (GPU if available)
x = x.to(device)

# Create directly on device
y = torch.zeros(3, 4, device=device)
```

**Key insight:** Both data and model must be on the **same device**.

## Device Management Pattern



Always move data to device **inside** the training loop.

## The Three Components

1. **Dataset**: Holds data, implements `__getitem__` and `__len__`
2. **Transforms**: Preprocessing and augmentation
3. **DataLoader**: Batching, shuffling, parallel loading

## Transform Composition

```
from torchvision import transforms

data_transforms = {
    'train': transforms.Compose([
        transforms.RandomResizedCrop(224),
        transforms.RandomHorizontalFlip(),
        transforms.ToTensor(),
        transforms.Normalize([0.485, 0.456, 0.406],
                           [0.229, 0.224, 0.225])
    ]),
    'val': transforms.Compose([
        transforms.Resize(256),
        transforms.CenterCrop(224),
        transforms.ToTensor(),
        transforms.Normalize([0.485, 0.456, 0.406],
                           [0.229, 0.224, 0.225])
    ])
}
```

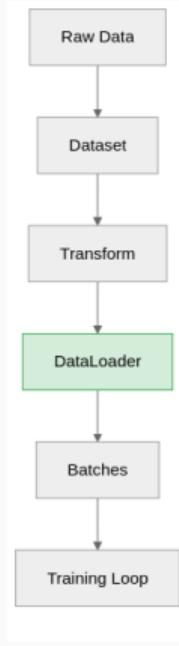
## DataLoader Configuration

```
from torch.utils.data import DataLoader

train_loader = DataLoader(
    dataset,
    batch_size=32,
    shuffle=True,          # Randomize order each epoch
    num_workers=4          # Parallel data loading
)
```

Parameter	Training	Validation
shuffle	True	False
batch_size	Tunable	Same or larger
drop_last	Often True	False

# Data Loading Flow



DataLoader handles batching and parallel loading automatically.

## The nn.Module Contract

Every PyTorch model inherits from nn.Module:

```
import torch.nn as nn

class MyNetwork(nn.Module):
    def __init__(self):
        super().__init__()
        # Define layers here

    def forward(self, x):
        # Define computation here
        return x
```

Two required methods: `__init__` and `forward`.

## Layer Definition Pattern

```
class SimpleNet(nn.Module):
    def __init__(self, input_dim, hidden_dim, num_classes):
        super().__init__()
        self.linear1 = nn.Linear(input_dim, hidden_dim)
        self.linear2 = nn.Linear(hidden_dim, num_classes)
        self.relu = nn.ReLU()

    def forward(self, x):
        x = self.relu(self.linear1(x))
        x = self.linear2(x)
        return x
```

Layers defined in `__init__`, used in `forward`.

## What nn.Module Gives You

- **Parameter tracking:** All weights automatically registered
- **Device movement:** `.to(device)` moves all parameters
- **State management:** `state_dict()` captures all weights
- **Mode switching:** `train()` and `eval()` methods
- **Gradient control:** Works with autograd system

## Common Layer Types

Layer	Purpose	Example
<code>nn.Linear</code>	Fully connected	<code>nn.Linear(784, 256)</code>
<code>nn.Conv2d</code>	2D convolution	<code>nn.Conv2d(3, 64, 3)</code>
<code>nn.ReLU</code>	Activation	<code>nn.ReLU()</code>
<code>nn.Dropout</code>	Regularization	<code>nn.Dropout(0.5)</code>
<code>nn.BatchNorm2d</code>	Normalization	<code>nn.BatchNorm2d(64)</code>

## The Core Pattern

```
for epoch in range(num_epochs):
    model.train()
    for inputs, labels in train_loader:
        inputs = inputs.to(device)
        labels = labels.to(device)

        optimizer.zero_grad()      # 1. Clear gradients
        outputs = model(inputs)    # 2. Forward pass
        loss = criterion(outputs, labels) # 3. Compute loss
        loss.backward()            # 4. Backward pass
        optimizer.step()          # 5. Update weights
```

This 5-step pattern is **universal** in PyTorch training.

## Why Zero Gradients?

- PyTorch **accumulates** gradients by default
- Without zeroing: gradients from previous batch add up
- Must call `optimizer.zero_grad()` before each backward pass

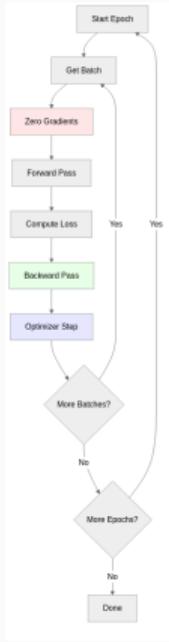
*# These are equivalent:*

```
optimizer.zero_grad()
```

*# or manually:*

```
for param in model.parameters():
    if param.grad is not None:
        param.grad.zero_()
```

# Training Loop Visualization



## Common Loss Functions

```
import torch.nn as nn

# Classification
criterion = nn.CrossEntropyLoss()    # Multi-class
criterion = nn.BCEWithLogitsLoss()   # Binary

# Regression
criterion = nn.MSELoss()           # Mean squared error
criterion = nn.L1Loss()            # Mean absolute error
```

**Note:** CrossEntropyLoss combines LogSoftmax + NLLLoss.

## Optimizer Selection

```
import torch.optim as optim

# Stochastic Gradient Descent
optimizer = optim.SGD(model.parameters(),
                      lr=0.01,
                      momentum=0.9)

# Adam (adaptive learning rate)
optimizer = optim.Adam(model.parameters(),
                       lr=0.001)
```

## Optimizer Comparison

Optimizer	Pros	Cons	When to Use
SGD	Simple, well-understood	Slow convergence	With LR scheduling
SGD+Momentum	Faster than SGD	Extra hyperparameter	CNNs often
Adam	Fast, adaptive	May generalize worse	Default starting point
AdamW	Better regularization	-	Transformers

## Passing Parameters to Optimizer

# All parameters

```
optimizer = optim.Adam(model.parameters(), lr=0.001)
```

# Only specific layers

```
optimizer = optim.Adam(model.fc.parameters(), lr=0.001)
```

# Different learning rates per layer

```
optimizer = optim.Adam([
    {'params': model.features.parameters(), 'lr': 1e-4},
    {'params': model.classifier.parameters(), 'lr': 1e-3}
])
```

## Two Model Modes

```
model.train()    # Training mode  
model.eval()    # Evaluation mode
```

These modes affect behavior of certain layers.

## Layers Affected by Mode

Layer	train()	eval()
Dropout	Randomly drops units	No dropout applied
BatchNorm	Uses batch statistics	Uses running statistics
Other layers	No change	No change

**Critical:** Always set correct mode before forward pass.

## Correct Usage Pattern

*# Training*

```
model.train()  
  
for inputs, labels in train_loader:  
    outputs = model(inputs)  
    loss = criterion(outputs, labels)  
    loss.backward()  
    optimizer.step()
```

*# Evaluation*

```
model.eval()  
  
with torch.no_grad():  
    for inputs, labels in val_loader:  
        outputs = model(inputs)  
        # Compute metrics...
```

## Common Mistake

```
# WRONG: Forgot to set eval mode
def predict(model, inputs):
    return model(inputs) # Dropout still active!
```

```
# CORRECT: Set eval mode
def predict(model, inputs):
    model.eval()
    with torch.no_grad():
        return model(inputs)
```

Dropout during inference = **non-deterministic** predictions.

## `torch.no_grad()`

Disables gradient computation:

```
with torch.no_grad():
    outputs = model(inputs)
    # No gradients computed
    # Cannot call loss.backward()
```

### Benefits:

- Reduces memory usage
- Speeds up computation
- Essential for inference

## `torch.set_grad_enabled()`

More flexible gradient control:

```
# Conditional gradient computation
with torch.set_grad_enabled(phase == 'train'):
    outputs = model(inputs)
    loss = criterion(outputs, labels)

    if phase == 'train':
        loss.backward()
        optimizer.step()
```

Useful for combined train/val loops.

## Memory Implications

Context	Gradients Stored	Memory	Use Case
Default	Yes	High	Training
no_grad	No	Low	Inference
inference_mode	No	Lowest	Production

```
# Lowest memory for inference
with torch.inference_mode():
    outputs = model(inputs)
```

## Gradient Context Flow



## The state\_dict Pattern

```
# Get all model parameters as dictionary
state = model.state_dict()

# Keys are parameter names
print(state.keys())
# ['linear1.weight', 'linear1.bias', 'linear2.weight', ...]
```

This is how PyTorch serializes models.

## Saving and Loading Models

```
# Save model weights
torch.save(model.state_dict(), 'model.pth')

# Load model weights
model = MyNetwork() # Create architecture first
model.load_state_dict(torch.load('model.pth'))
model.eval()
```

**Best practice:** Save state\_dict, not entire model.

## Saving Complete Checkpoints

```
# Save everything for resuming training
checkpoint = {
    'epoch': epoch,
    'model_state_dict': model.state_dict(),
    'optimizer_state_dict': optimizer.state_dict(),
    'loss': loss,
}
torch.save(checkpoint, 'checkpoint.pth')

# Resume training
checkpoint = torch.load('checkpoint.pth')
model.load_state_dict(checkpoint['model_state_dict'])
optimizer.load_state_dict(checkpoint['optimizer_state_dict'])
```

## Copying Model Weights

```
import copy

# Deep copy for best model tracking
best_model_wts = copy.deepcopy(model.state_dict())

# Later, restore best weights
model.load_state_dict(best_model_wts)
```

Useful for early stopping and model selection.

## The `requires_grad` Flag

Every tensor has a `requires_grad` attribute:

```
# Check if parameter needs gradients
for name, param in model.named_parameters():
    print(name, param.requires_grad)

# Freeze a parameter
param.requires_grad = False
```

Frozen parameters receive **no gradient updates**.

## Freezing for Transfer Learning

```
# Load pretrained model
model = models.resnet18(pretrained=True)

# Freeze all parameters
for param in model.parameters():
    param.requires_grad = False

# Replace and train only final layer
model.fc = nn.Linear(512, num_classes)
# New layer has requires_grad=True by default
```

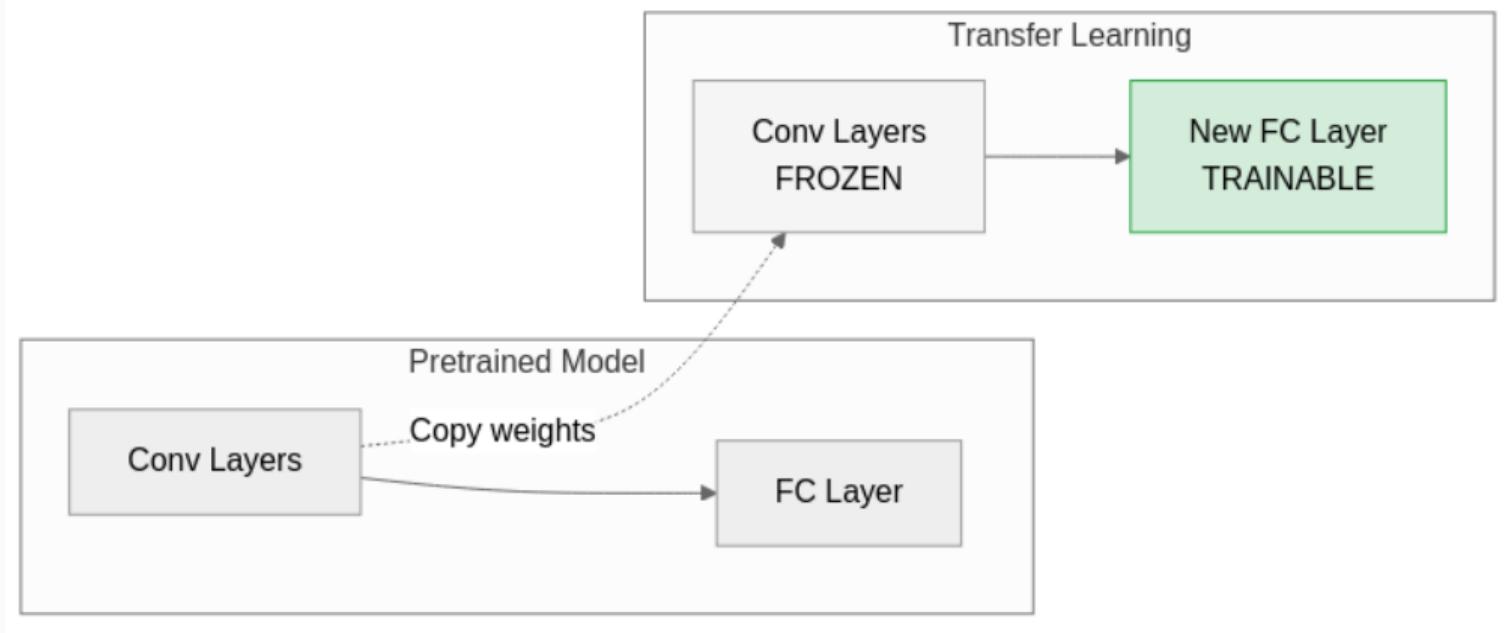
Only the new layer receives gradients.

## Selective Freezing

```
# Freeze by name pattern
for name, param in model.named_parameters():
    if 'layer1' in name or 'layer2' in name:
        param.requires_grad = False

# Only optimize unfrozen parameters
optimizer = optim.Adam(
    filter(lambda p: p.requires_grad, model.parameters()),
    lr=0.001
)
```

## Freezing Visualization



## Why Schedule Learning Rates?

- **Large LR early:** Fast initial progress
- **Small LR later:** Fine-tune near optimum
- Helps escape local minima
- Improves final convergence

## Common Schedulers

```
from torch.optim.lr_scheduler import StepLR, ExponentialLR

# Decay by factor every N epochs
scheduler = StepLR(optimizer, step_size=7, gamma=0.1)

# Decay by factor every epoch
scheduler = ExponentialLR(optimizer, gamma=0.95)

# Cosine annealing
scheduler = CosineAnnealingLR(optimizer, T_max=100)
```

## Scheduler in Training Loop

```
for epoch in range(num_epochs):
    model.train()
    for inputs, labels in train_loader:
        optimizer.zero_grad()
        outputs = model(inputs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

# Step scheduler AFTER epoch
scheduler.step()

print(f"LR: {scheduler.get_last_lr()}")
```

**Note:** Call `scheduler.step()` after each epoch.

## Scheduler Comparison

Scheduler	Behavior	Use Case
StepLR	Discrete drops	Most common
ExponentialLR	Smooth decay	Gradual refinement
CosineAnnealingLR	Cosine curve	Vision models
ReduceLROnPlateau	Adaptive	When loss plateaus

## Loading Pretrained Models

```
from torchvision import models

# Load with pretrained weights
model = models.resnet18(pretrained=True)

# Or with new API (PyTorch 2.0+)
model = models.resnet18(
    weights=models.ResNet18_Weights.IMAGENET1K_V1
)
```

Models pretrained on ImageNet's 1000 classes.

## Modifying for Your Task

```
# Original final layer
print(model.fc)
# Linear(in_features=512, out_features=1000)

# Replace for binary classification
num_classes = 2
model.fc = nn.Linear(model.fc.in_features, num_classes)

# Move to device
model = model.to(device)
```

New layer has random weights, ready to train.

## Two Transfer Learning Strategies

### Feature Extraction (freeze backbone):

```
for param in model.parameters():
    param.requires_grad = False
model.fc = nn.Linear(512, num_classes) # Only this trains
```

### Fine-tuning (train everything):

```
model.fc = nn.Linear(512, num_classes)
# All parameters trainable, but use small LR
optimizer = optim.SGD(model.parameters(), lr=0.001)
```

## When to Use Each Strategy

---

Scenario	Strategy	Learning Rate
Small dataset, similar domain	Feature extraction	Normal
Large dataset	Fine-tuning	Small for pretrained layers
Different domain	Fine-tuning	Very small
Very small dataset	Feature extraction + simple head	Normal

---

## What is Parameterization?

A way to add transformations to existing weights:

```
import torch.nn.utils.parametrize as parametrize

class ScaleWeight(nn.Module):
    def forward(self, W):
        return W * 2.0 # Double all weights

parametrize.register_parametrization(
    layer, "weight", ScaleWeight()
)
```

The weight is now computed through the parameterization.

## Use Cases

- **Spectral normalization:** Constrain weight norms
- **Weight orthogonalization:** Maintain orthogonal weights
- **LoRA:** Add low-rank updates (see LoRA lecture)
- **Pruning:** Mask out weights

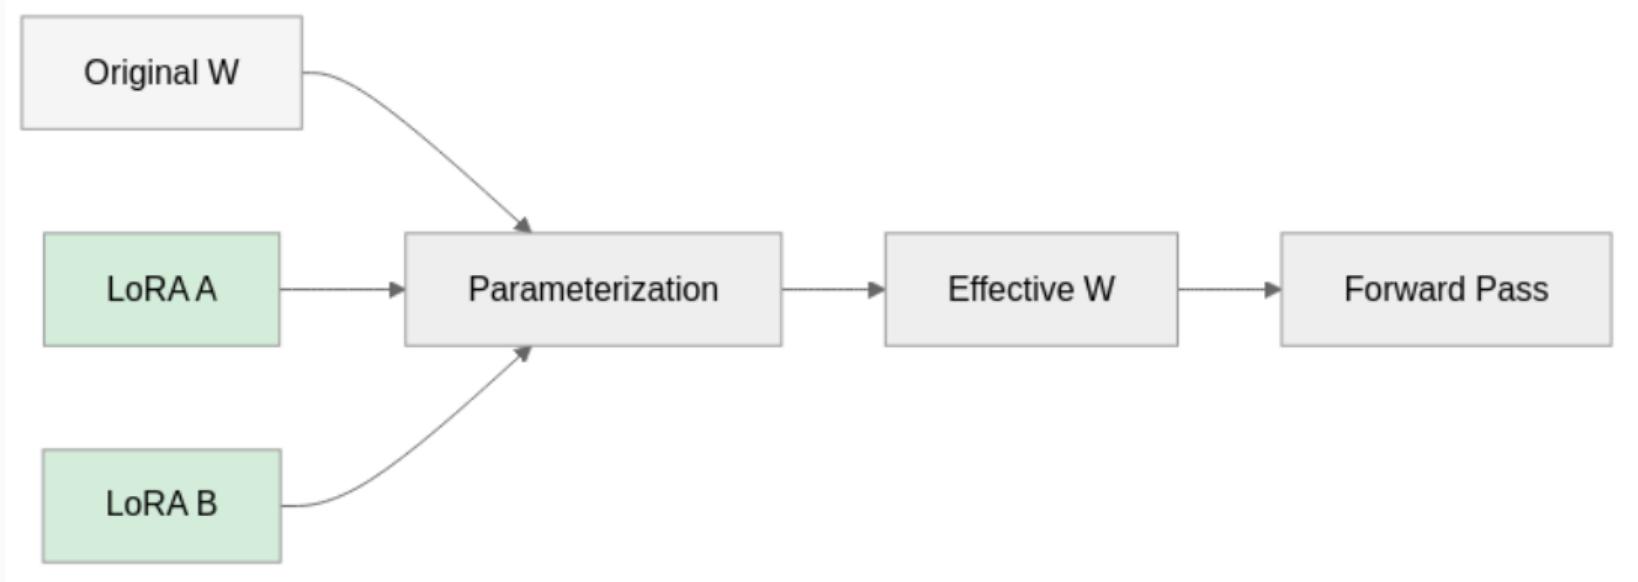
## LoRA as Parameterization

```
class LoRAParametrization(nn.Module):
    def __init__(self, features_in, features_out, rank=8):
        super().__init__()
        self.lora_A = nn.Parameter(torch.randn(rank, features_out))
        self.lora_B = nn.Parameter(torch.zeros(features_in, rank))

    def forward(self, W):
        return W + self.lora_B @ self.lora_A
```

Original weights + low-rank update.

## Parameterization Flow



## Complete Training Script Template

```
import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader

# 1. Device setup
device = torch.device("cuda" if torch.cuda.is_available()
                      else "cpu")

# 2. Data loading
train_loader = DataLoader(train_dataset, batch_size=32,
                           shuffle=True)
val_loader = DataLoader(val_dataset, batch_size=32)
```

## Template (continued)

### # 3. Model setup

```
model = MyNetwork().to(device)
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=0.001)
scheduler = optim.lr_scheduler.StepLR(optimizer,
                                      step_size=7, gamma=0.1)
```

### # 4. Training loop

```
for epoch in range(num_epochs):
    model.train()
    for inputs, labels in train_loader:
        inputs, labels = inputs.to(device), labels.to(device)

        optimizer.zero_grad()
        outputs = model(inputs)
        loss = criterion(outputs, labels)
```

## Template (continued)

```
# 5. Validation
model.eval()
with torch.no_grad():
    for inputs, labels in val_loader:
        inputs, labels = inputs.to(device), labels.to(device)
        outputs = model(inputs)
        # Compute metrics...
    scheduler.step()

# 6. Save model
torch.save(model.state_dict(), 'final_model.pth')
```

## M03 Notebooks Overview

---

Notebook	Patterns Demonstrated
ConvolutionalNet_Classifier	Data transforms, pretrained models, training loop
Overfitting_Regularization	Model comparison, regularization
Transfer_Learning_LoRA	Parameter freezing, parameterization
TSNE_EMBEDDING	Feature extraction from trained models

---

We will explore these patterns hands-on in the live demos.

## Core Patterns

1. **Device Management:** Always define device, move data and model
2. **Data Pipeline:** Dataset → Transforms → DataLoader
3. **nn.Module:** `__init__` defines layers, `forward` defines computation
4. **Training Loop:** `zero_grad` → `forward` → `loss` → `backward` → `step`
5. **Loss & Optimizer:** Choose based on task and model

## Advanced Patterns

6. **Train/Eval Mode:** Set correctly before inference
7. **Gradient Contexts:** Use `no_grad()` for inference
8. **State Management:** Save/load with `state_dict()`
9. **Parameter Freezing:** Control `requires_grad` for transfer learning
10. **LR Scheduling:** Decay learning rate during training
11. **Pretrained Models:** Load, modify, fine-tune
12. **Parameterization:** Add transformations to weights

## What to Remember

- PyTorch patterns are **consistent** across projects
- The training loop structure is **universal**
- Mode switching (train/eval) is **critical** for correct behavior
- State management enables **reproducibility** and **transfer learning**

## PyTorch Documentation

- PyTorch Tutorials: [pytorch.org/tutorials](https://pytorch.org/tutorials)
- nn.Module documentation: [pytorch.org/docs/stable/nn.html](https://pytorch.org/docs/stable/nn.html)
- Parametrization: [pytorch.org/tutorials/intermediate/parametrizations.html](https://pytorch.org/tutorials/intermediate/parametrizations.html)

## Further Reading

- D2L Book: [d2l.ai](https://d2l.ai/) - Dive into Deep Learning
- CS231n: [cs231n.github.io](https://cs231n.github.io) - CNN for Visual Recognition
- Transfer Learning Tutorial: [pytorch.org/tutorials/beginner/transfer\\_learning\\_tutorial.html](https://pytorch.org/tutorials/beginner/transfer_learning_tutorial.html)